



RICH AI

Mukund Iyer, Nico Van den Hooff,
Rakesh Pandey, Shiva Jena

MDS Mentor: Arman Seyed-Ahmadi



Agenda

- 1) Background and scope (Mukund)
- 2) Data (Mukund)
- 3) Machine learning models (Shiv)
- 4) Point clouds (Mukund)
- 5) Deep learning (Nico)
- 6) Model summary (Nico)
- 7) Data product (Shiv)



Background and scope

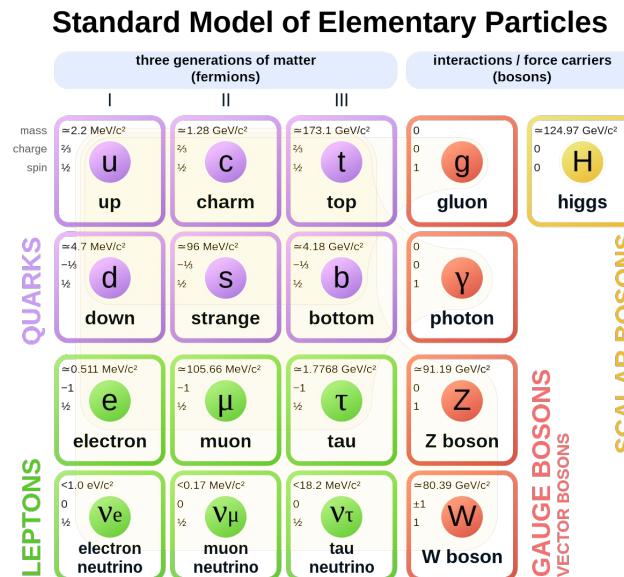


The Standard Model describes the composition of matter

Physicists' Objective: verify the “Standard Model”



Source: <https://exoplanets.nasa.gov/what-is-an-exoplanet/what-is-the-universe/>



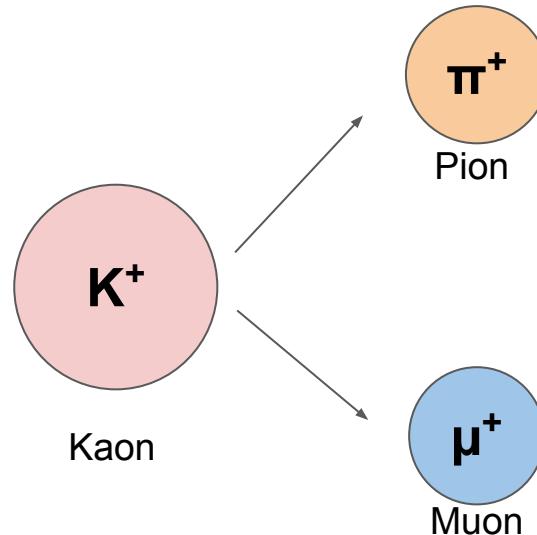
The NA62 experiment at CERN is used to study pions



Source: <https://home.cern/science/experiments/na62>

Geneva, Switzerland

Decay:



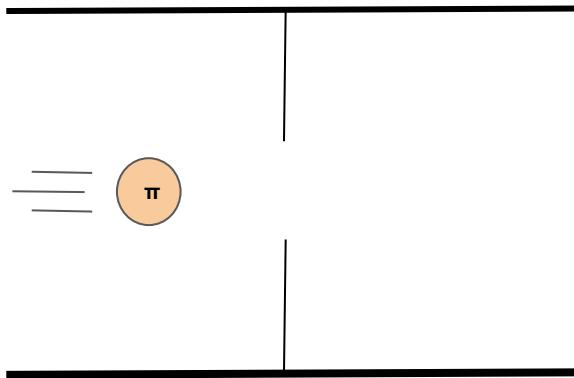
**Need to distinguish
between these particles!**

Need to analyze 2 sets of features to determine the subatomic particle



*cross sectional view

1. Particle Motion

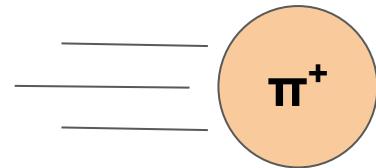


Need to analyze 2 sets of features to determine the subatomic particle

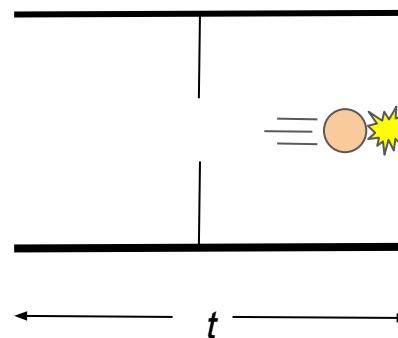


*cross sectional view

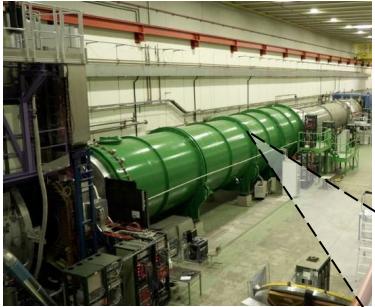
Particle
Momentum



Particle movement time
(CHOD time)

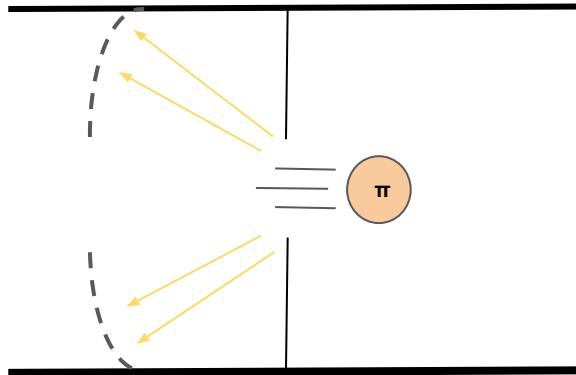


Need to analyze 2 sets of features to determine the subatomic particle

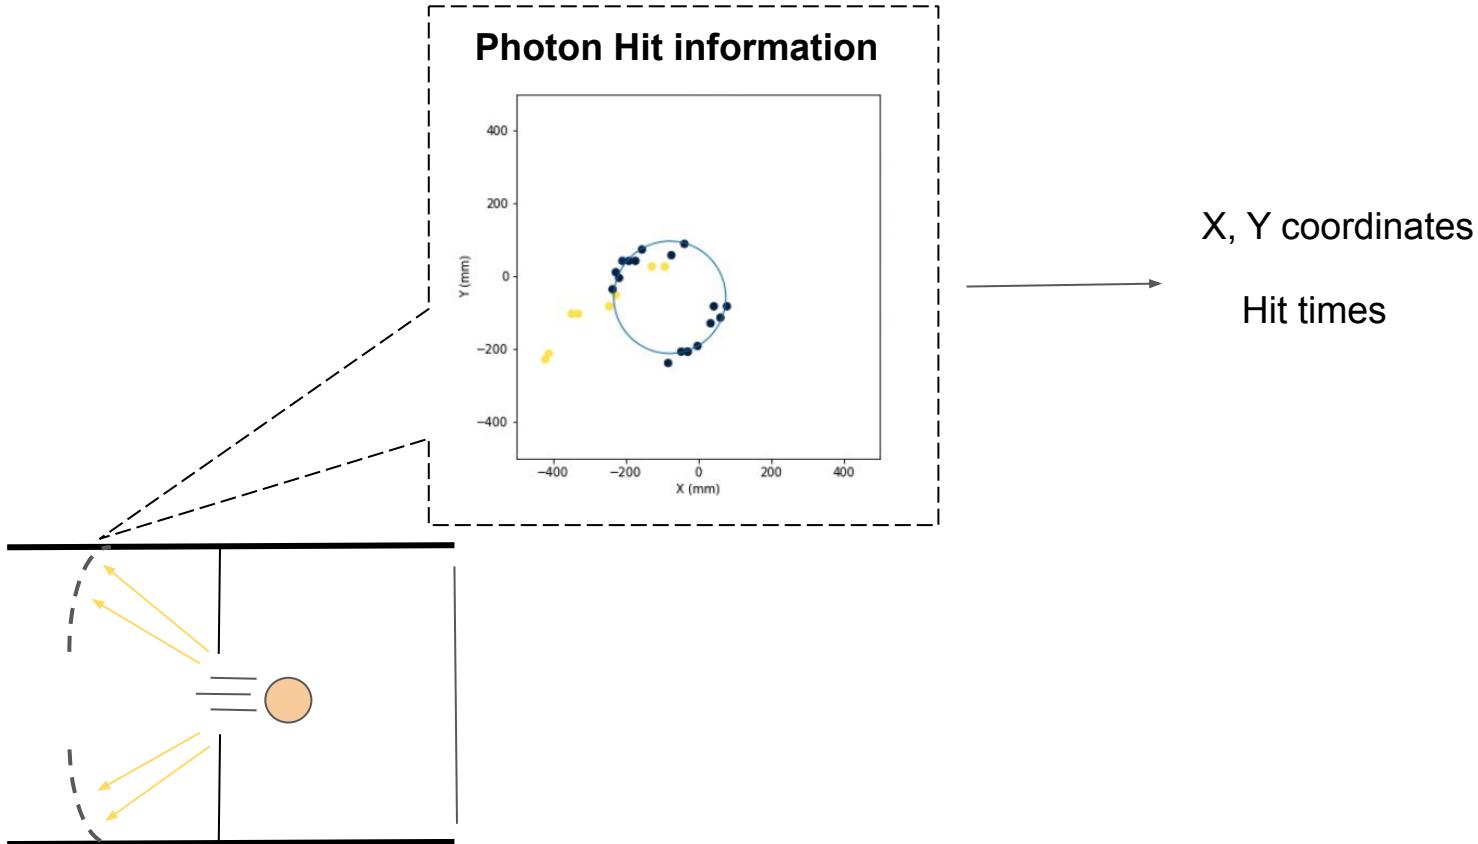


*cross sectional view

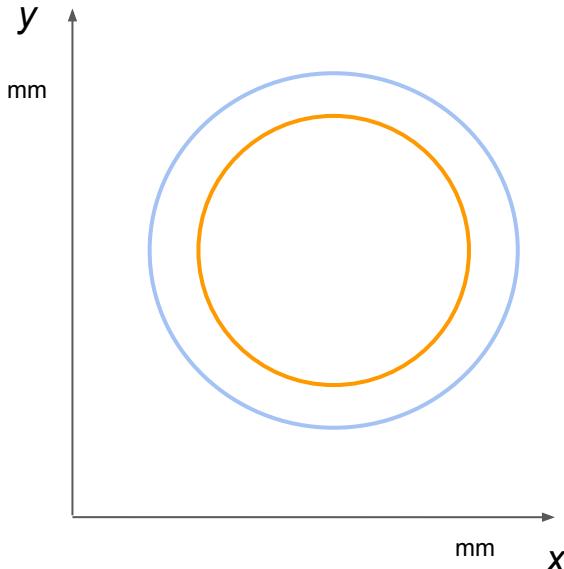
2. Photon (light) hits



Multiple features are extracted from the photon hits for each event

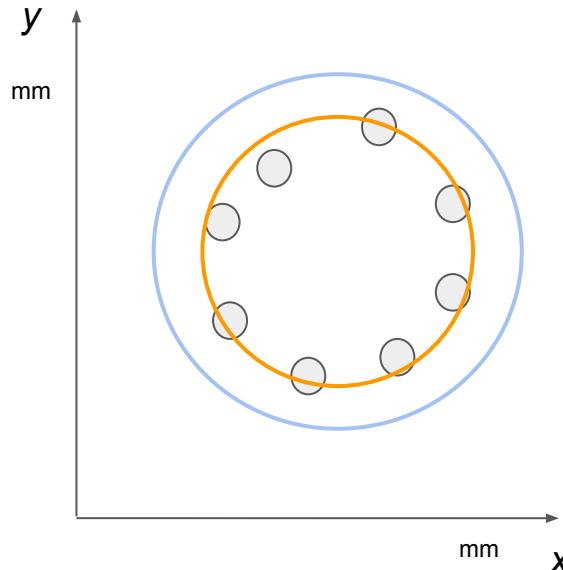


Can distinguish particle by ring size... binary classification!



For a given momentum,
the ring size for a
pion is **smaller**
than for a **muon!**

TRIUMF use a MLE estimate to determine the particle



$P(\text{pion}) = 0.8$

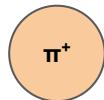
$P(\text{muon}) = 0.2$

This is a **pion!**

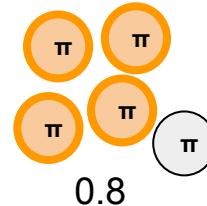
This method is purely analytical!

Objective: surpass the current performance of the MLE algorithm

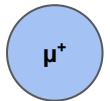
The two main metrics of interest are:



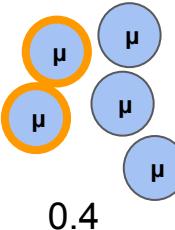
$$\text{Pion Efficiency} = \frac{\# \text{ Pions Identified}}{\# \text{ Total pions}}$$



Pion Recall!

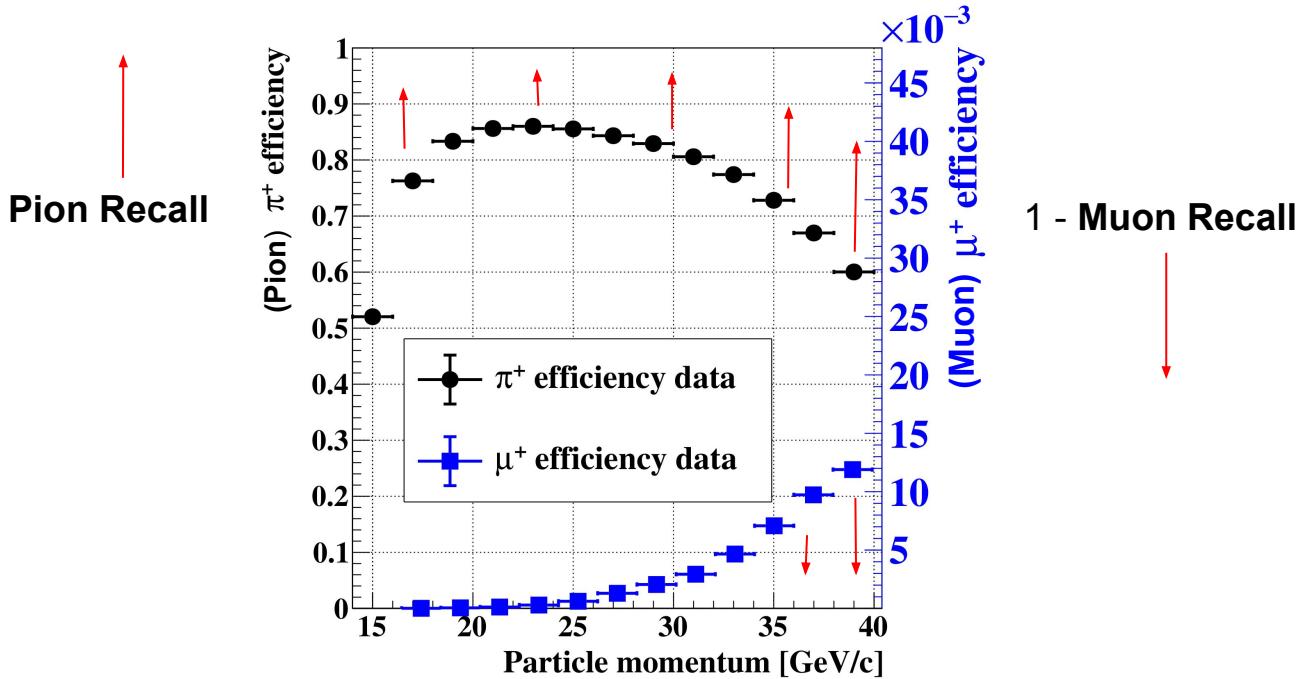


$$\text{Muon Efficiency} = \frac{\# \text{ Muons as pions}}{\# \text{ Total Muons}}$$

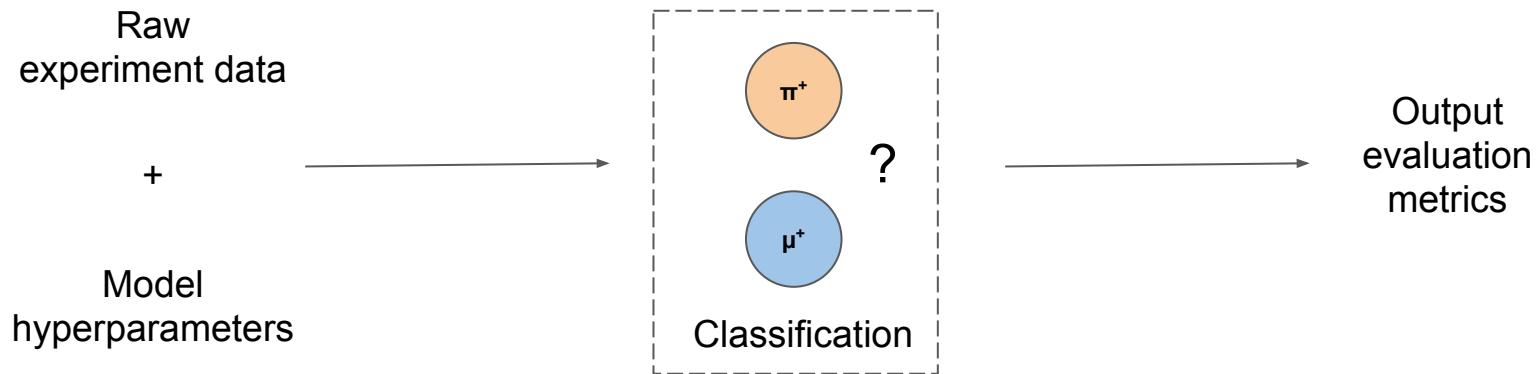


1 - Muon Recall!

Need to surpass the current performance generated from the MLE



Objective: build a modularized ML pipeline for a binary classifier





Data

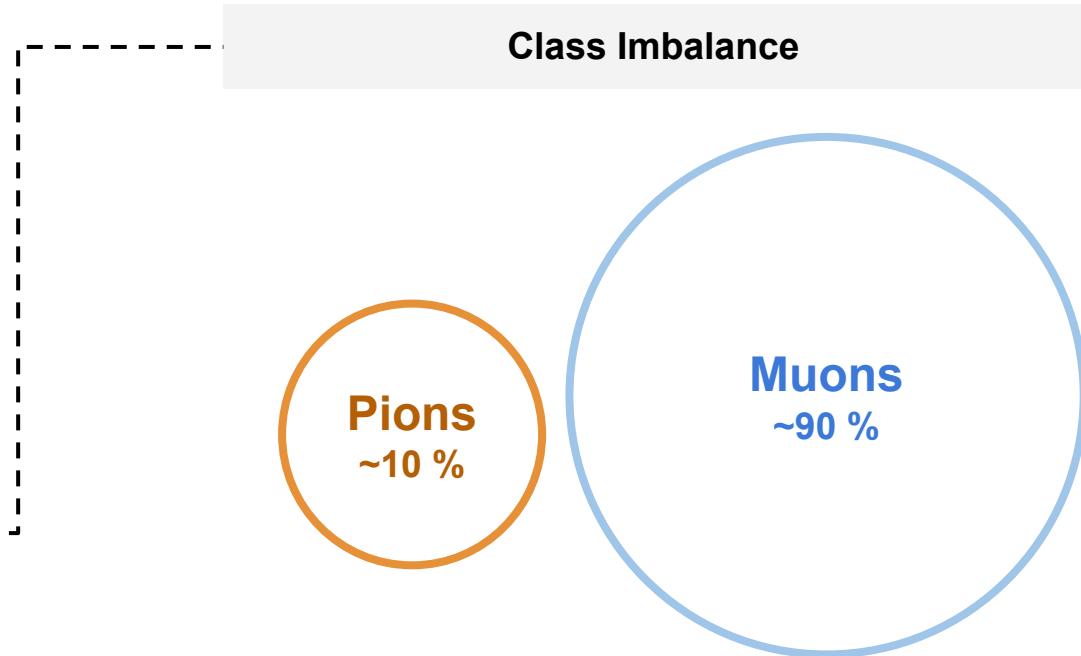


Data on 11.5 million events with a class imbalance

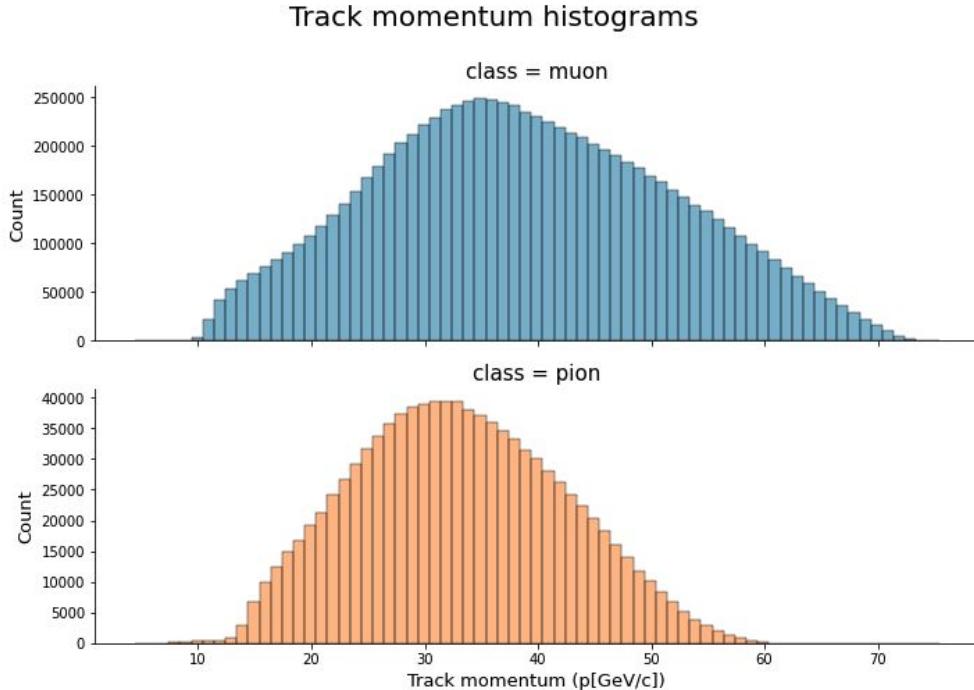
1 decay event = 1 example

Labelled data

~11.5 million examples



Systematic difference in the momentum distribution among particles



- Artifact of the experimental set-up
- **Do not** want ML models to pick up on this difference

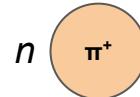
Undersample using momentum bins!

Procedure:

Slice data by
bins of size
 $10 \text{ GeV}/c^2$

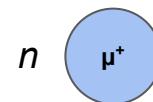


Count
number of
pions in bin

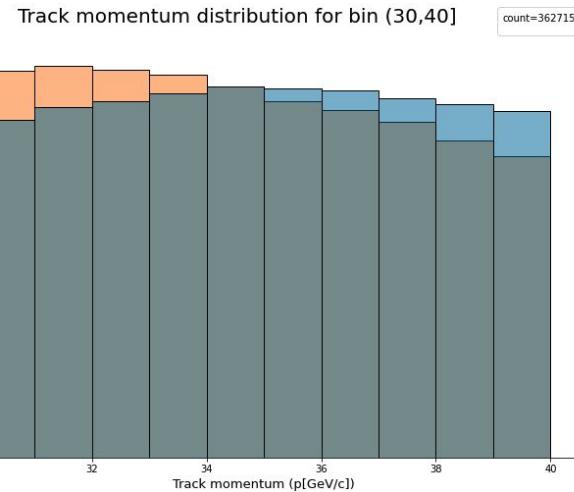
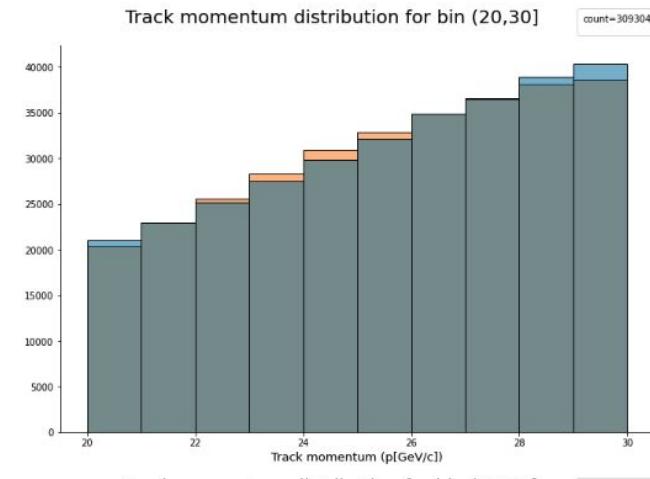


Sample equal
number of
muons

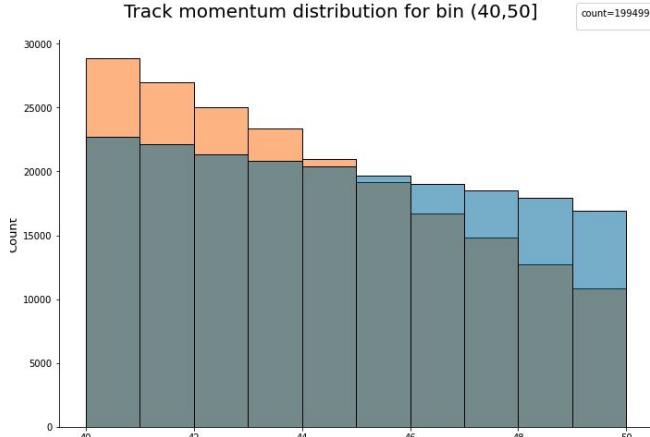
$\sim 2 \text{ million samples}$



Undersample using momentum bins!



Pion
Muon



- Distribution looks **even** across all bins
- Evidence that momentum is **controlled**



Baseline machine learning model
XGBoost



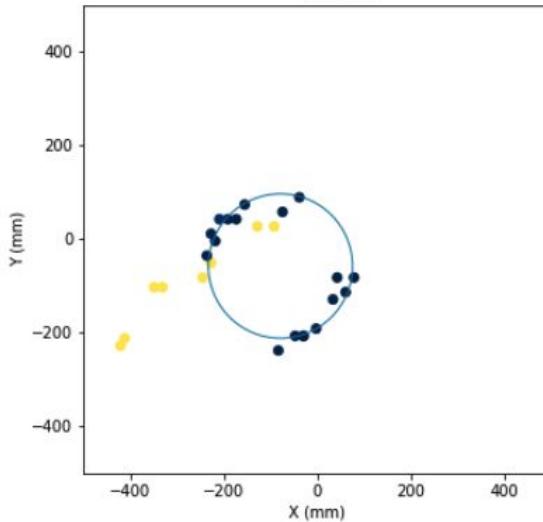
Benchmarking with XGBoost as baseline: Features

Features

1. MLE radius
2. # Hits
3. Momentum

Refresher: MLE radius and # Hits engineered from photon hit information

Photon Hit information



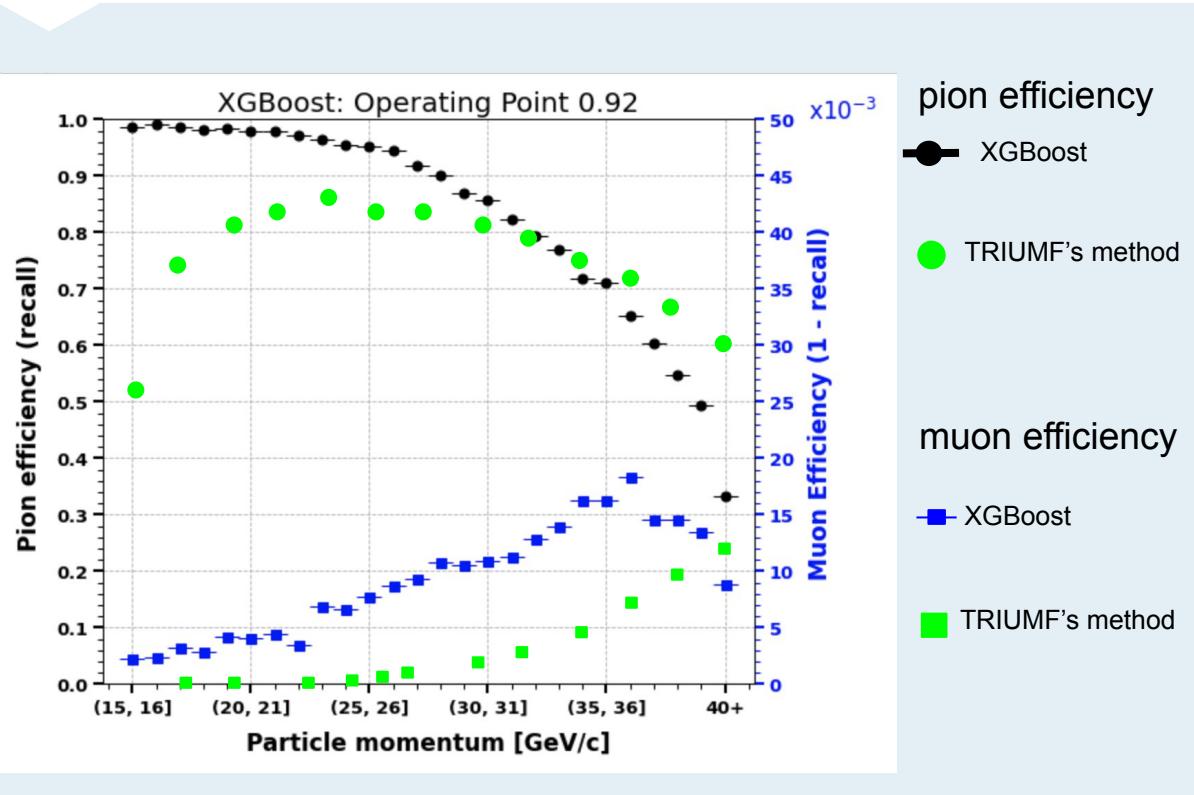
Ring Radius - MLE fit

- done by TRIUMF
- does not use X,Y coordinates

Hits - noise reduction

Benchmarking with XGBoost as baseline: Performance

Model performance against TRIUMF's method



Hyperparameter optimisation ?

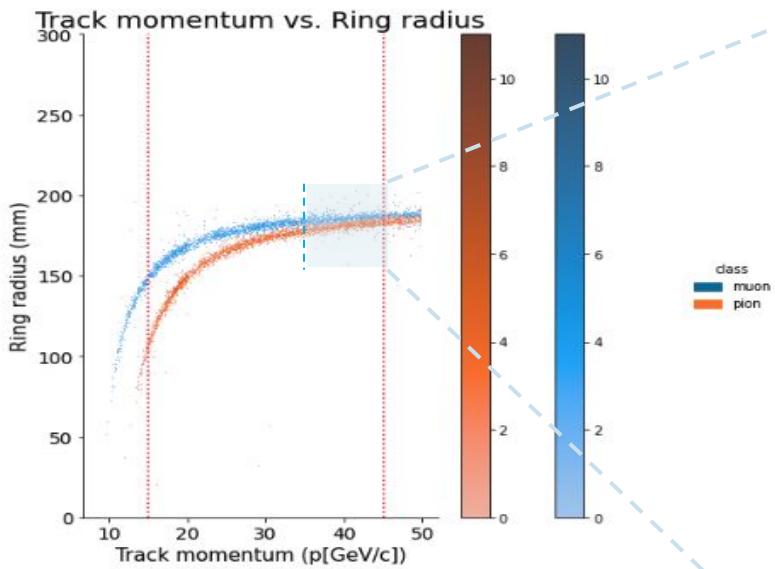
- Results did not improve significantly.

Why more complex models?

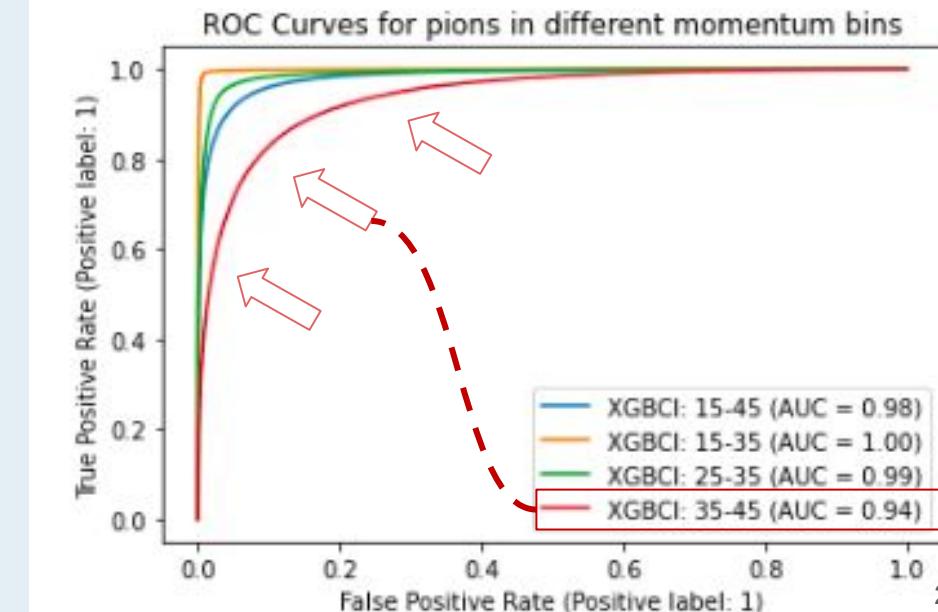
- Neural networks: **more accurate** in feature extraction (from photon hits data)

Benchmarking with XGBoost as baseline: Precise target area

Refresher: relationship between radius & momentum gets tricky in higher momentum



Discovery of a precise avenue
Better performance possible in higher momentum bins





Deep Learning
Point clouds



Point clouds are a set of points in space

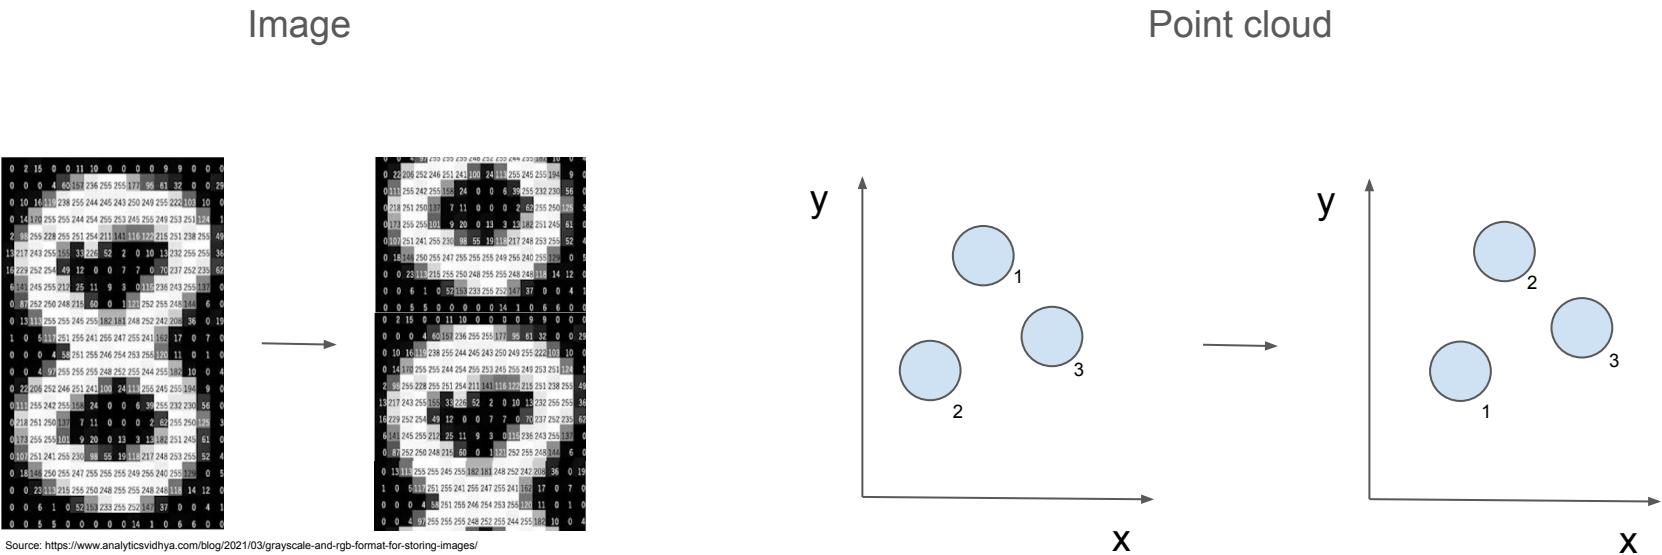


Source: <https://github.com/nikitakaraevv/pointnet>

Individual point contains **spatial** information

Together the points form an image or object

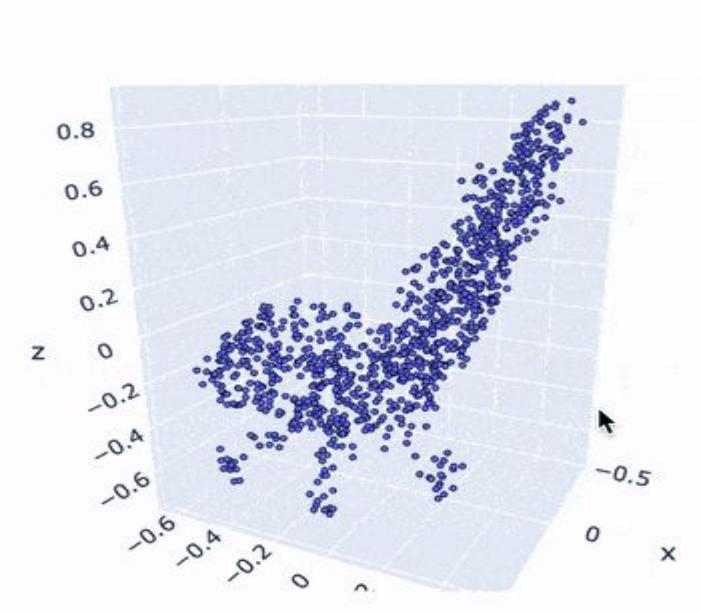
Point clouds do not have order



Changing the points changes the overall information

No change in information when changing points around

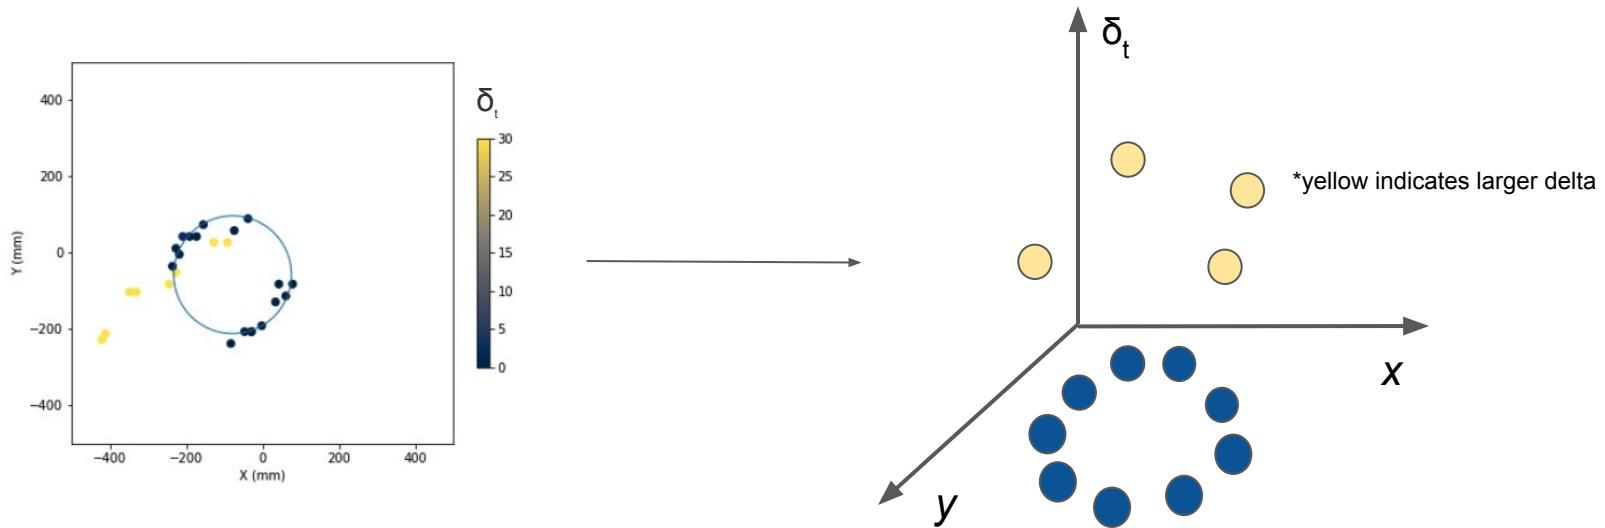
Point clouds are invariant to geometric transformations



Source: <https://github.com/nikitakaraevv/pointnet>

The object is still the same regardless of how it is rotated(or shift axis)!

Third dimension of time added to the hits data to generate Point clouds



$\bar{\delta}_t$ = difference between particle travel time and time of the photon hits ≈ 0



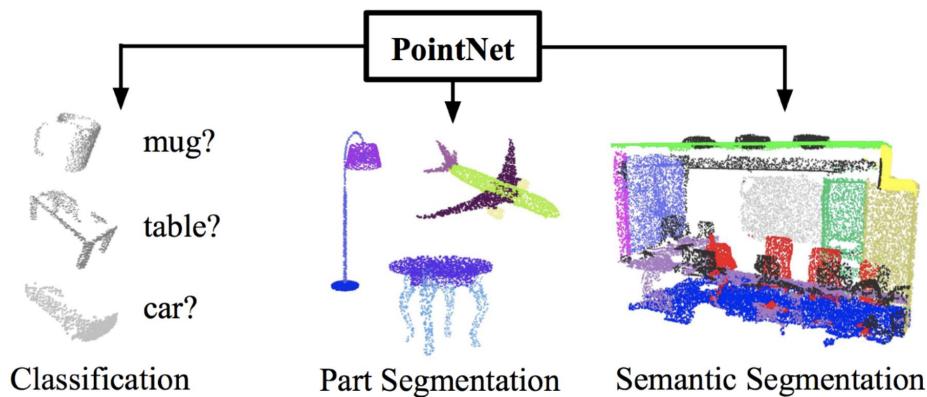
Deep Learning



The models that we used

Model 1: PointNet

- Developed by researchers at Stanford
- One of the first neural networks specifically meant for point clouds



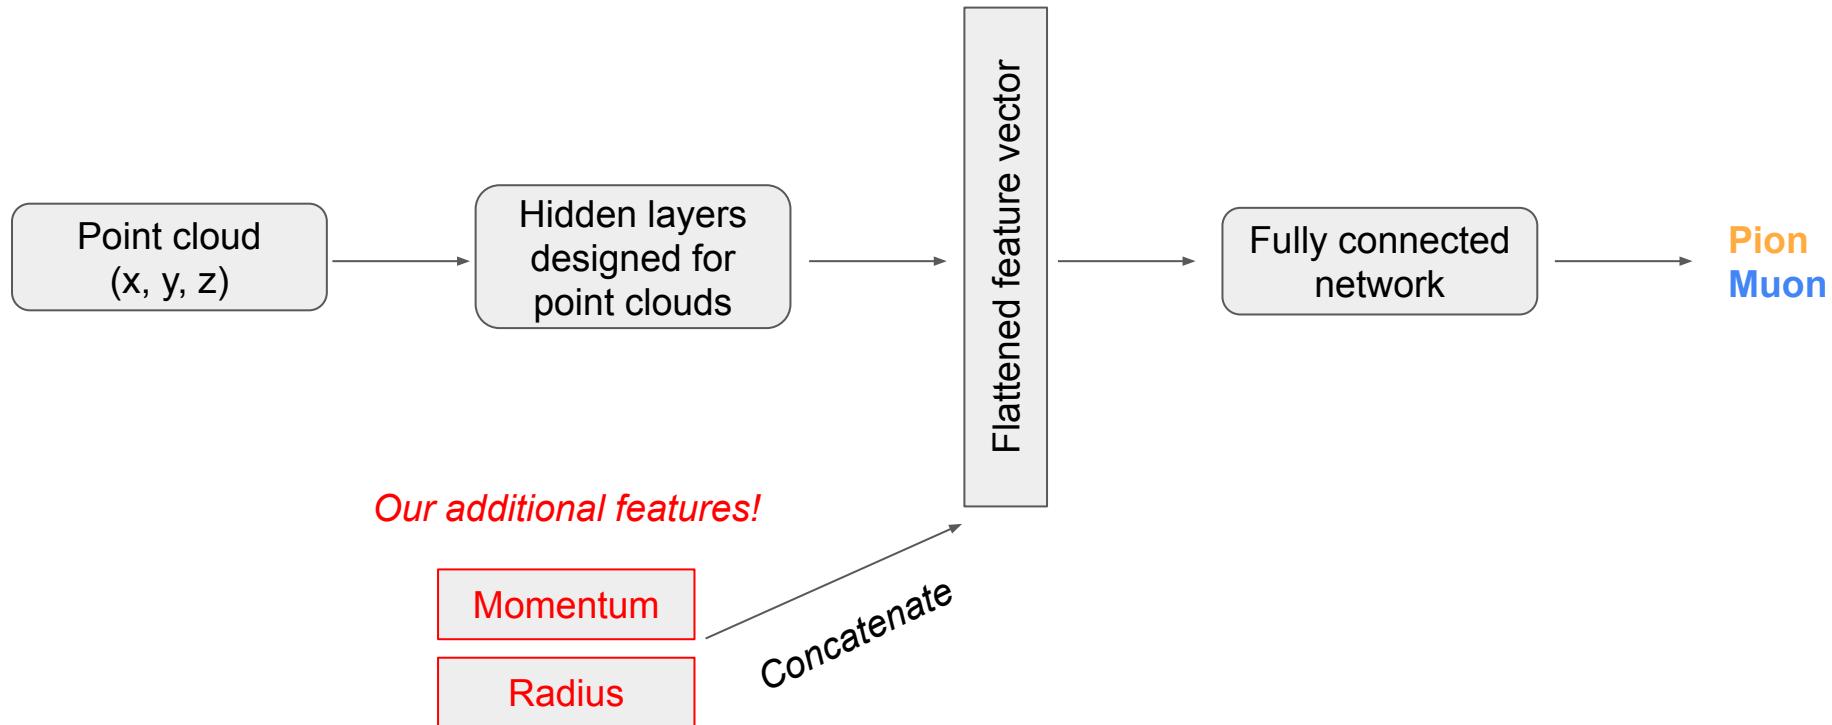
Model 2: Dynamic Graph CNN

- Developed by researchers at MIT
- Also intended for point cloud data
- Able to extract local information within pointcloud (PointNet cannot)

Dynamic Graph CNN Key Idea

“Edge convolutions”

High level model architecture (both models)





Deep Learning
Model #1: PointNet



How PointNet addresses problems with point clouds

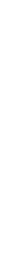
Problem 1

Point clouds are unordered sets of coordinates, changing the order within the set should not change the prediction



Problem 2

Rotating the point cloud should not change the corresponding prediction
(robust to spatial variability)



Solution 1

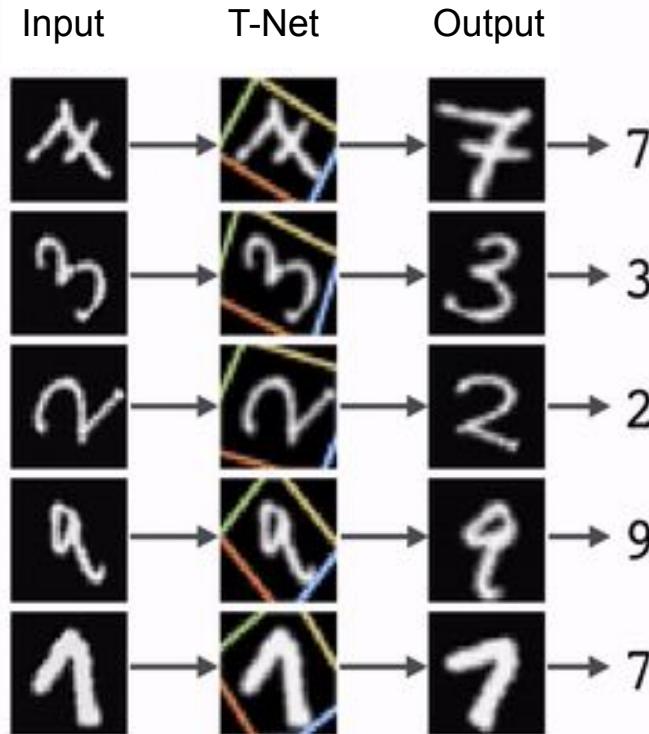
Apply max pooling (symmetric function) to transformed inputs

Solution 2

Spatial transformer network (“T-net”)

Spatial Transformer Network

- T-nets come from separate research
- T-nets are robust to input rotations
- PointNet uses a T-net to take advantage of this!





Deep Learning
PointNet Results



Tuning PointNet

Feature combinations

- Hits (x, y, z), momentum and radius
- Hits (x, y, z), momentum

Time filter values (delta)

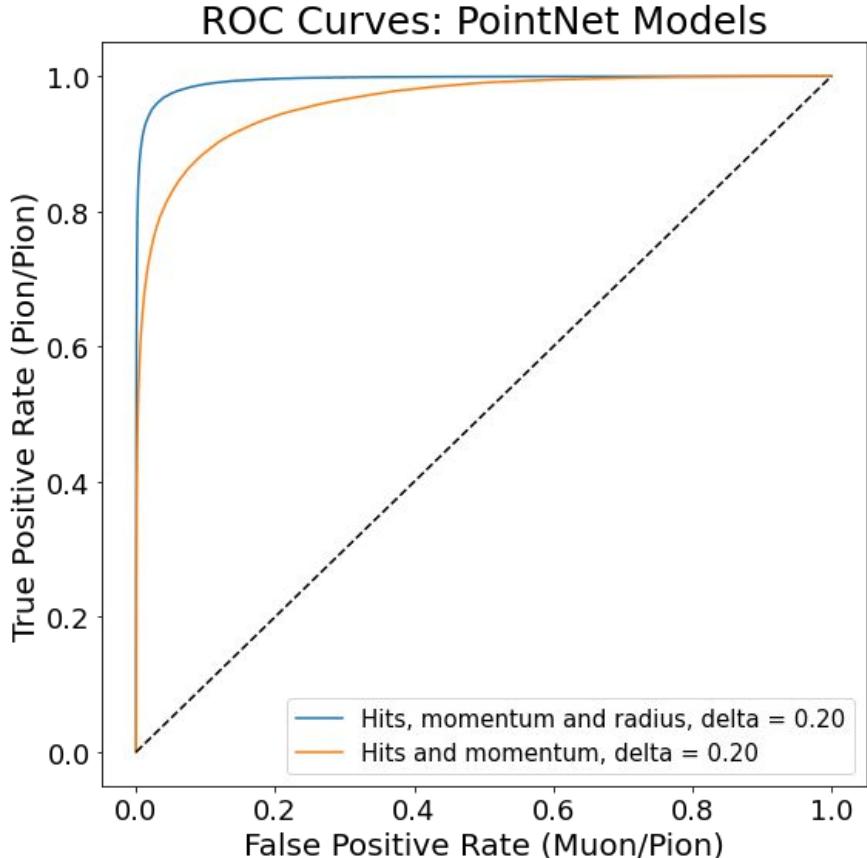
- 0.20ns to 0.50ns

Learning rate

- Constant learning rate
- Learning rate scheduler

Epochs

- Up to 24 epochs



Best PointNet model

Best feature combination

- Hits point cloud (x, y, z)
- Momentum
- Radius

Best hyperparameters

- Time delta = 0.20 ns
- 16 epochs
- Learning rate scheduler
 - 0.003 to 0.03 to 0.003

Chosen operating point (ROC)

- 0.93

Results

- Exceeds TRIUMFS prior pion efficiency results!
- Good muon efficiency, better than TRIUMF at *some* momentums

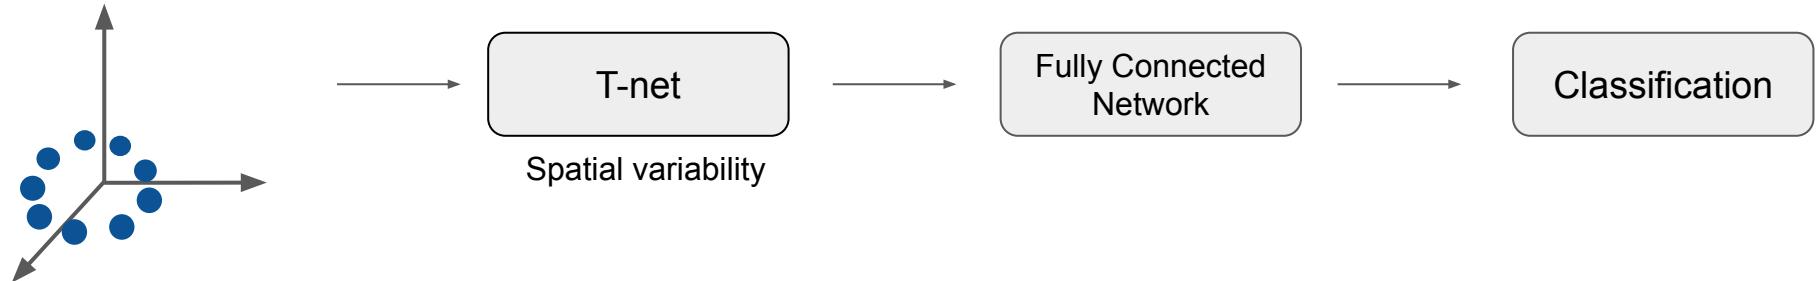


Deep Learning
Model #2: Dynamic Graph CNN

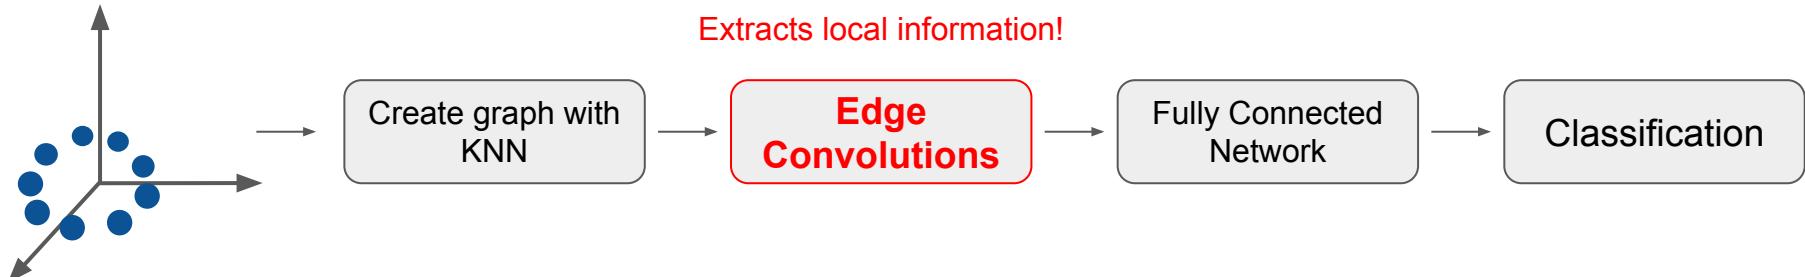


Comparing differences in model architectures

PointNet



Dynamic Graph CNN





Deep Learning
Dynamic Graph CNN Results



Tuning Dynamic Graph CNN

Feature combinations

- Hits (x, y, z), momentum and radius
- Hits (x, y, z), momentum

k for KNN graph

- 8 to 20 nearest neighbors

Time filter values (delta)

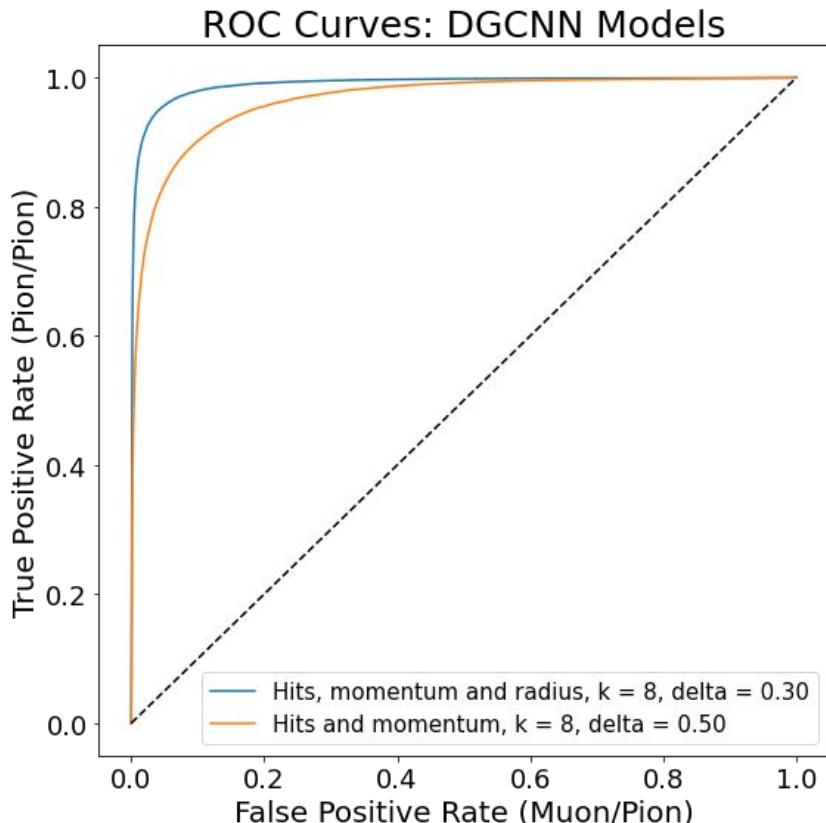
- 0.20ns to 0.50ns

Learning rate

- Constant learning rate
- Learning rate scheduler

Epochs

- Up to 12 epochs



Dynamic Graph CNN best model summary

Best feature combination

- Hits point cloud (x, y, z)
- Momentum
- Radius

Best hyperparameters

- $k = 8$ nearest neighbors
- Time delta = 0.30 ns
- 8 epochs
- Learning rate scheduler
 - 0.003 to 0.03 to 0.003

Chosen operating point (ROC)

- 0.96

Results

- *Similar* pion efficiency to PointNet
- Worse muon efficiency to PointNet



Deep learning
Overall summary of model results



Deep learning overall model results

Model performance

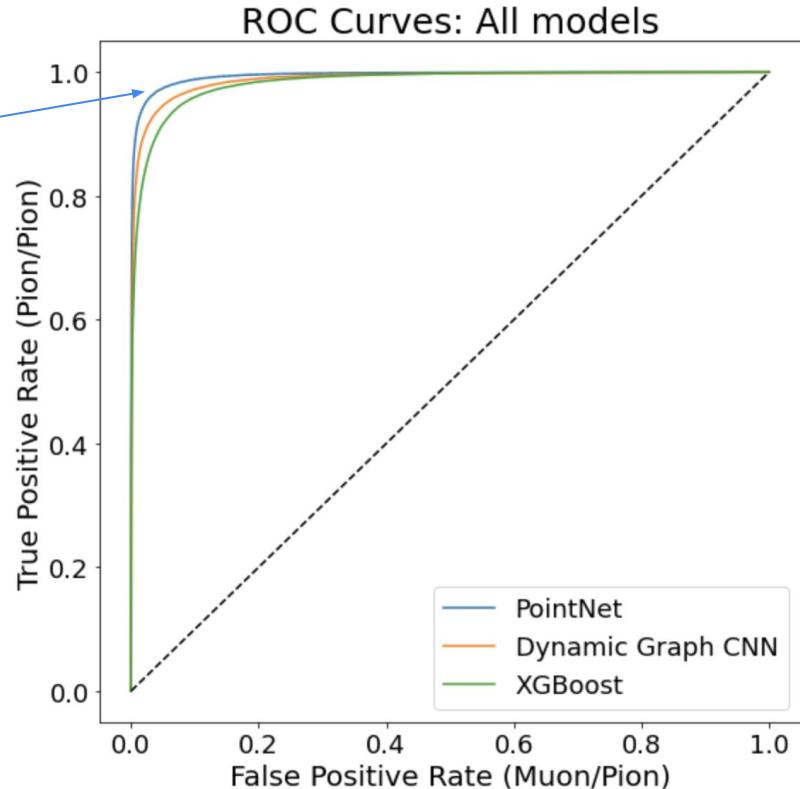
- Overall best model: PointNet
- Runner up: DGCNN
- Both models beat baseline!

PointNet Pros

- Exceeds all prior TRIUMF pion efficiency
- Pions more accurately identified
- Maintains strong muon efficiency

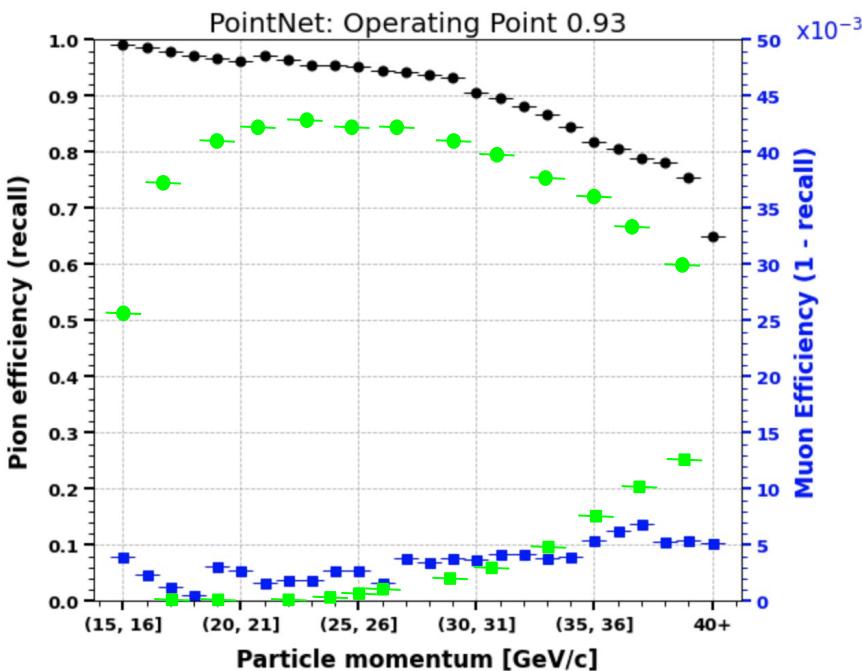
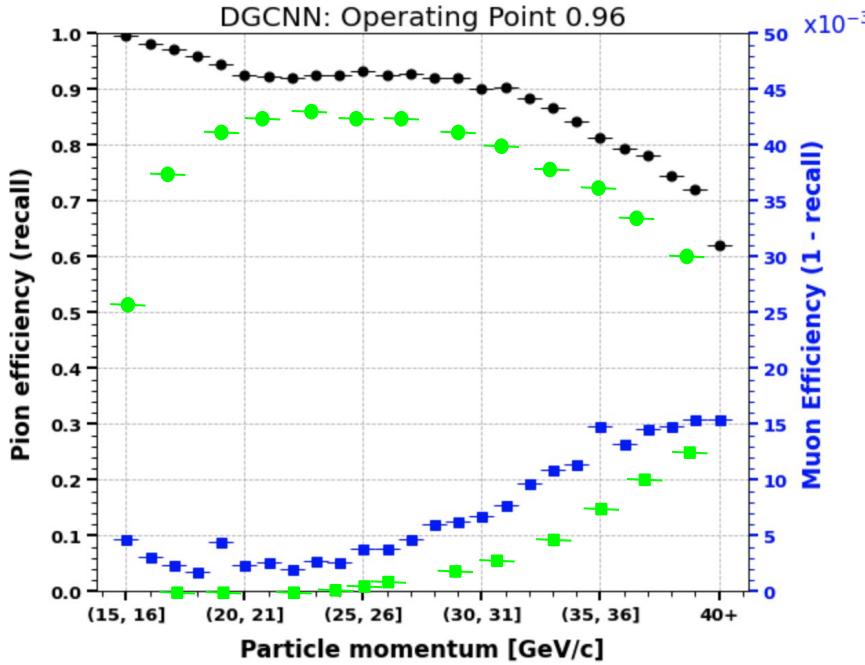
PointNet Cons

- Longer training time (double DGCNN)
- More parameters, potential to overfit



DGCNN vs. PointNet vs. TRIUMF

Prior TRIUMF performance



Better pion efficiency while *maintaining* good muon efficiency!

Potential suggestions for the future

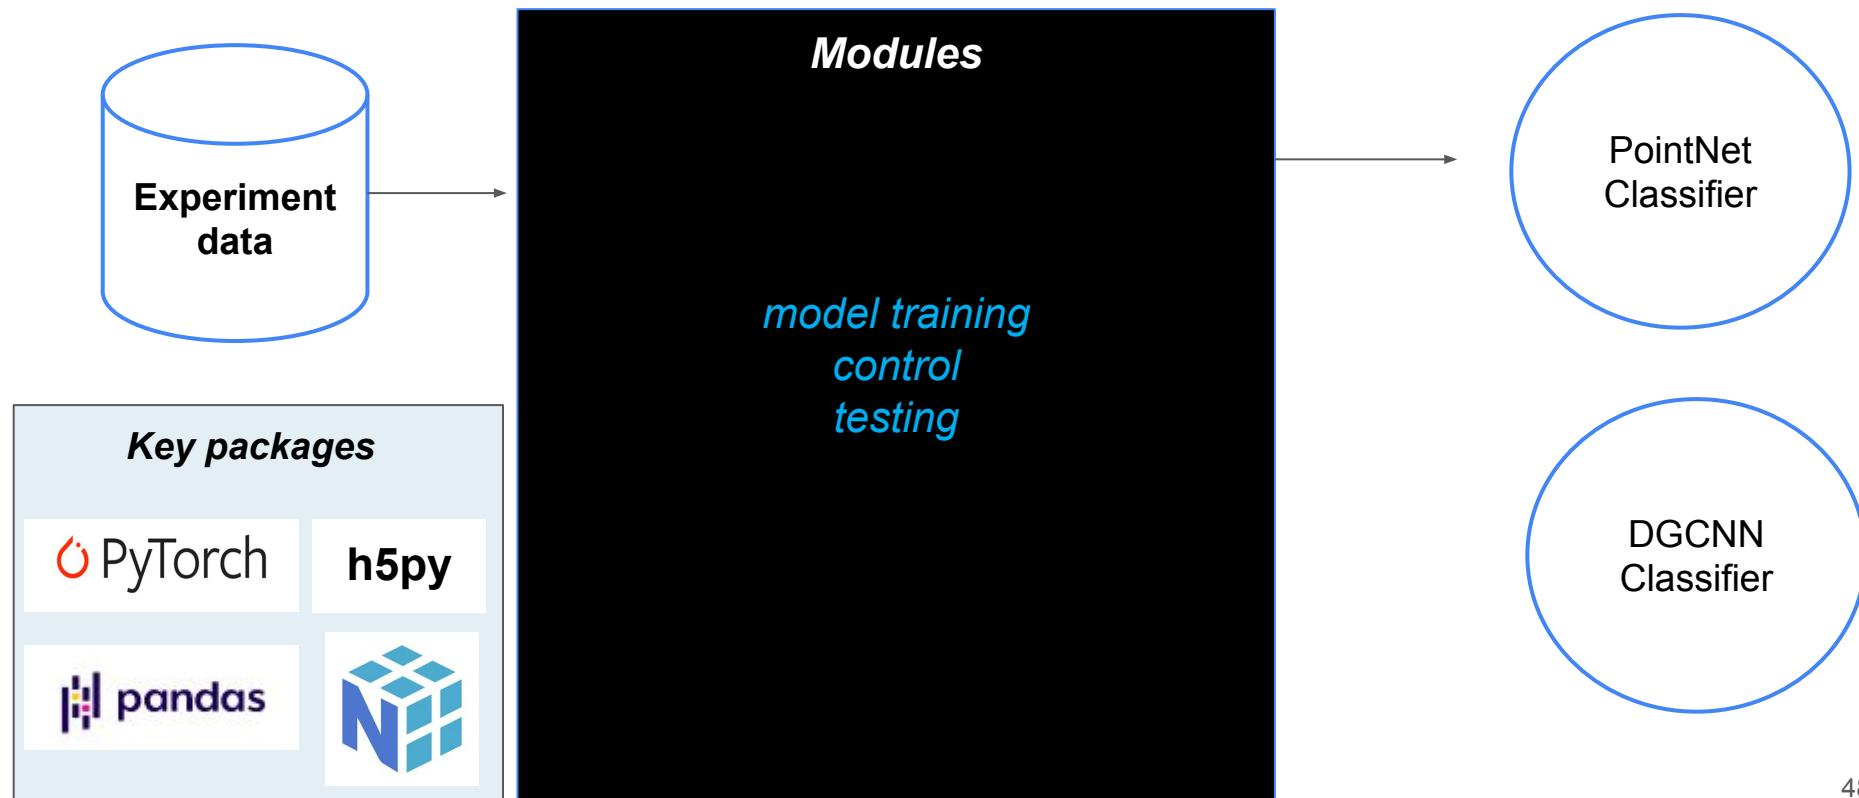
- 1) Further tuning of **hyperparameters** (e.g. Bayesian optimization)
- 2) Try **other models** (e.g. PointNet++)
- 3) Add **more muons** to the training dataset to help the model identify muons more accurately and further reduce muon efficiency

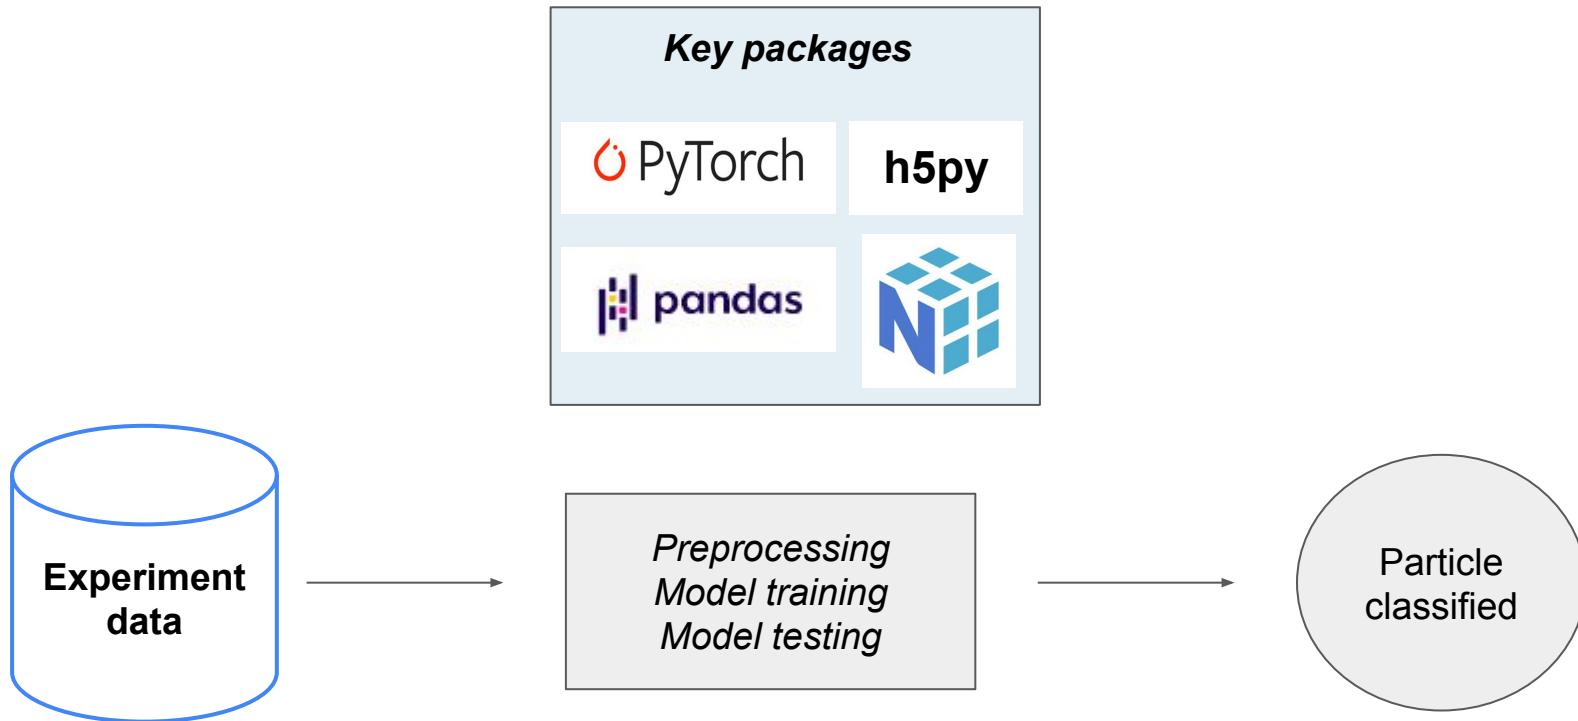


Final data product



Final data product: Modularized Neural Network classifiers





Final data product: Modularized Neural Network classifiers

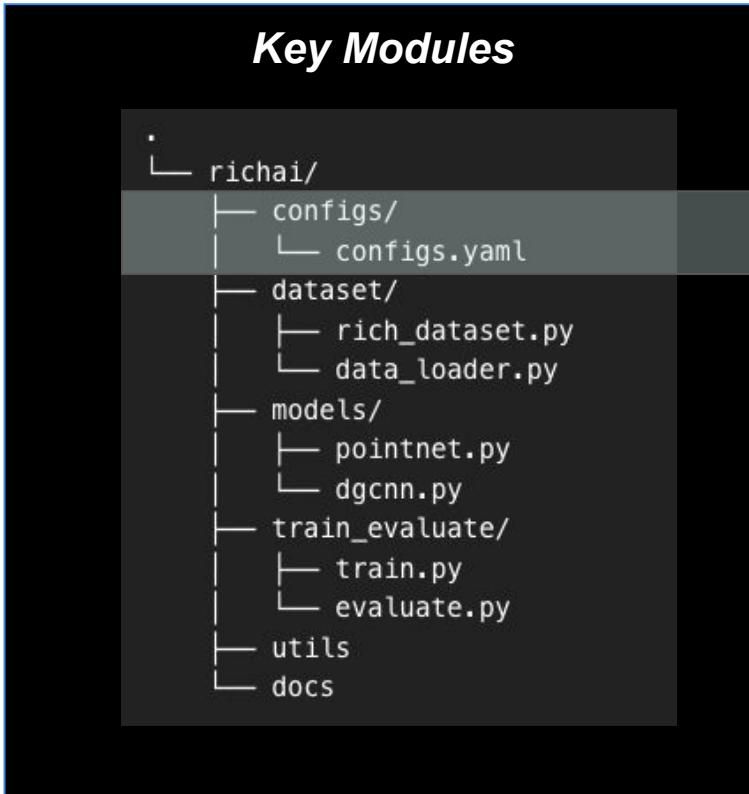


Key packages

PyTorch

h5py

pandas



Brief snapshot

```
data_loader:  
batch_size: 64  
num_workers: 8  
drop_last: True  
  
model:  
pointnet:  
    num_classes: 1  
    epochs: 16  
    saved_model: "saved_models/pointnet_1"  
    predictions: "saved_models/pointnet_1"  
    momentum: True  
  
gpu:  
    - 1  
    - 2  
    - 3  
  
seed: 2022
```

Final data product: Modularized Neural Network classifiers

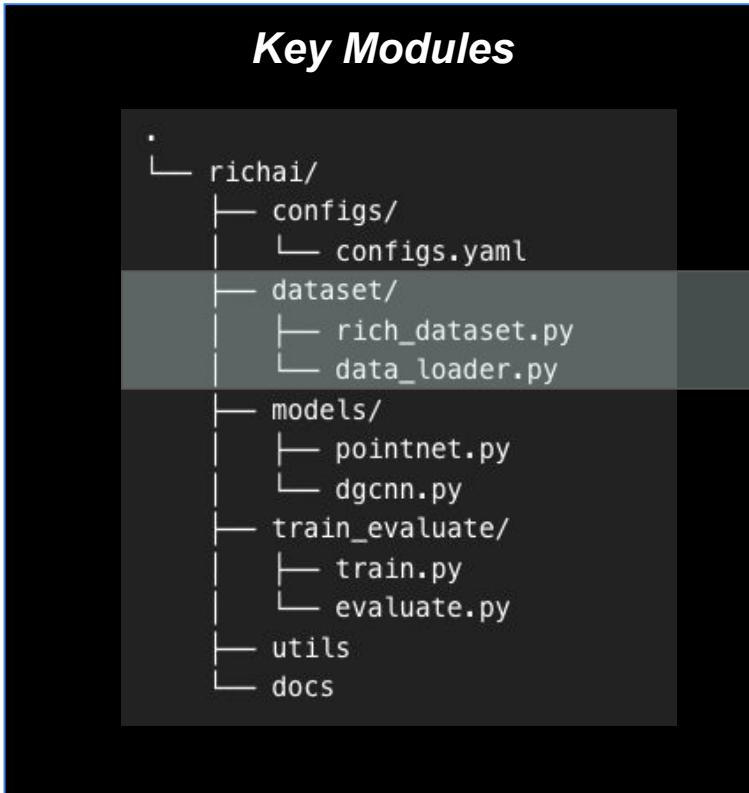


Key packages

PyTorch

h5py

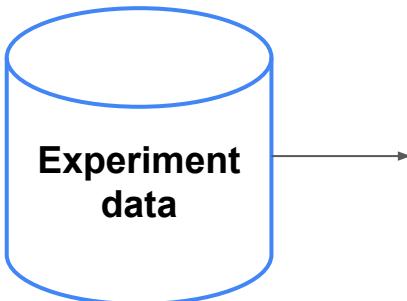
pandas



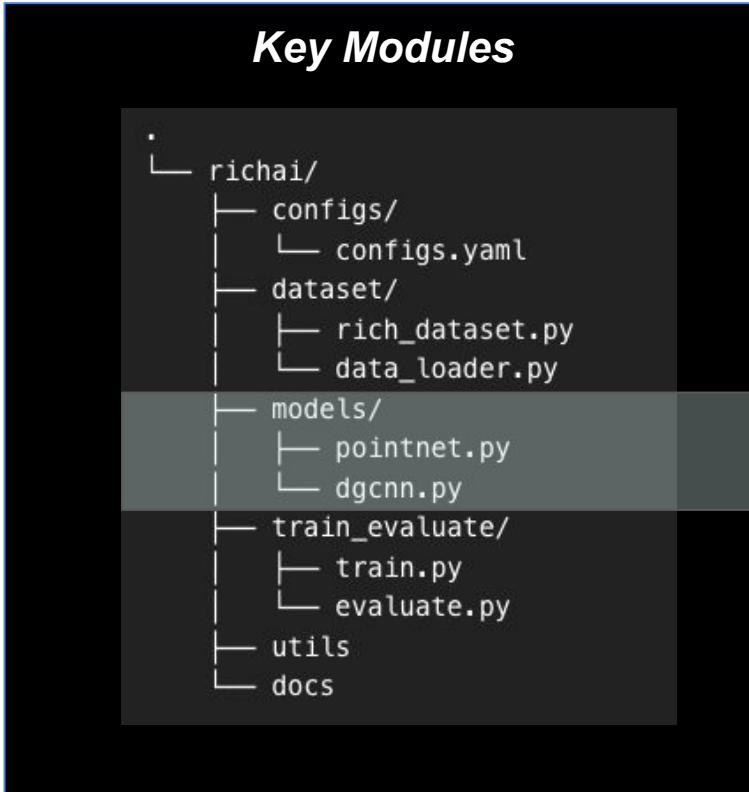
Brief snapshot

```
from torch.utils.data import DataLoader  
from torch.utils.data.sampler import SubsetRandomSampler  
from utils.helpers import get_config  
  
def data_loader(dset):  
    """Create pytorch data loader"""\n    train_loader = DataLoader(  
        dset,  
        batch_size=get_config("data_loader.batch_size"),  
        shuffle=False,  
        sampler=SubsetRandomSampler(dset.train_indices),  
        num_workers=get_config("data_loader.num_workers"),  
        drop_last=get_config("data_loader.drop_last"),  
    )  
  
    val_loader = DataLoader(  
        dset,  
        batch_size=get_config("data_loader.batch_size"),  
        shuffle=False,  
        sampler=SubsetRandomSampler(dset.val_indices),  
        num_workers=get_config("data_loader.num_workers"),  
        drop_last=get_config("data_loader.drop_last"),  
    )
```

Final data product: Modularized Neural Network classifiers



Key packages



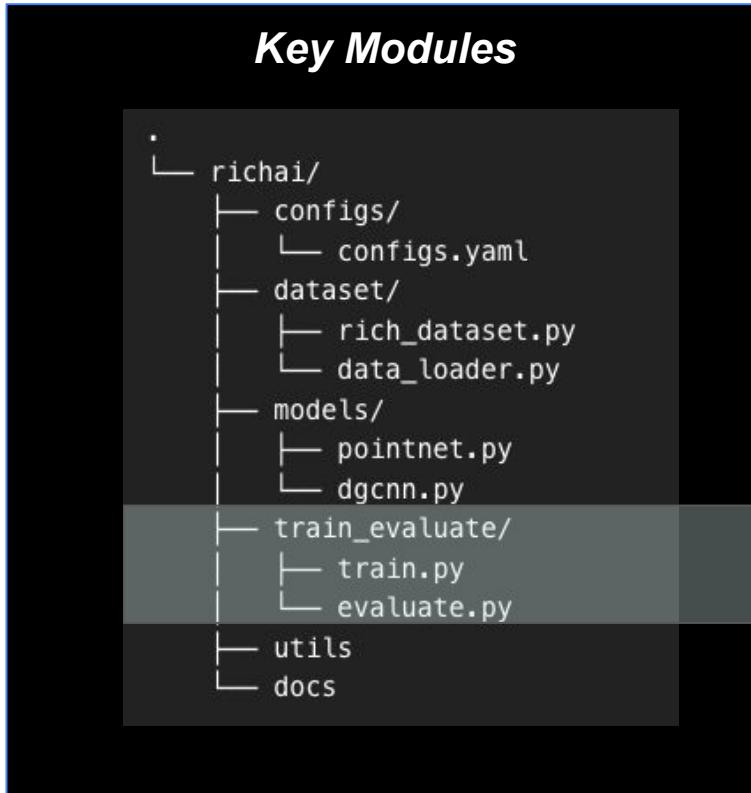
Brief snapshot

```
class STNkd(nn.Module):
    def __init__(self, k=64):
        super().__init__()
        self.k = k
        self.conv1 = nn.Conv1d(k, 64, 1)
        self.conv2 = nn.Conv1d(64, 128, 1)
        self.conv3 = nn.Conv1d(128, 1024, 1)
        self.fc1 = nn.Linear(1024, 512)
        self.fc2 = nn.Linear[512, 256]
        self.fc3 = nn.Linear[256, k * k]
        self.relu = nn.ReLU()

        self.bn1 = nn.BatchNorm1d(64)
        self.bn2 = nn.BatchNorm1d(128)
        self.bn3 = nn.BatchNorm1d(1024)
        self.bn4 = nn.BatchNorm1d(512)
        self.bn5 = nn.BatchNorm1d(256)

        self.iden = Variable(
            torch.from_numpy(np.eye(self.k).fla
).view(1, -1)
```

Final data product: Modularized Neural Network classifiers

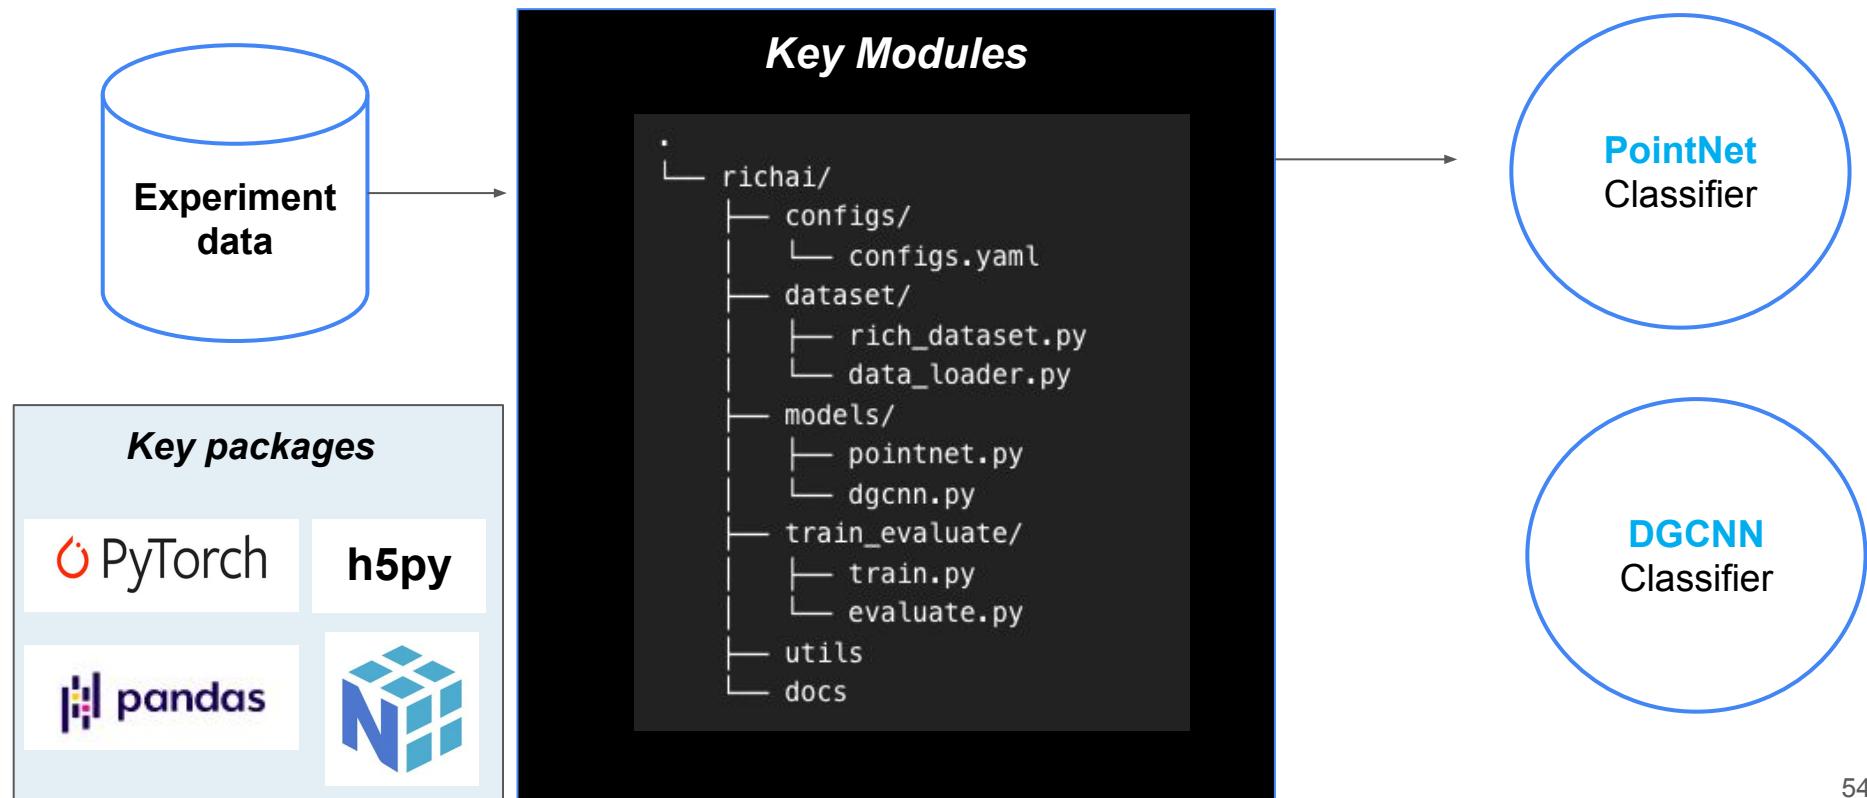


Brief snapshot

```
def trainer(  
    model,  
    optimizer,  
    train_loader,  
    val_loader=None,  
    criterion=None,  
    epochs=5,  
    scheduler=None,  
    device='cuda',  
    results=False,  
):  
    """Trainer for PyTorch"""\n\n    logger.info(f'Starting training...')  
    training_start = time.time()  
  
    train_losses, train_accs = [], []  
    valid_losses, valid_accs = [], []  
  
    model.to(device)
```

A dark grey rectangular area containing a snippet of Python code. The code defines a function 'trainer' that takes several parameters: 'model', 'optimizer', 'train_loader', 'val_loader' (set to None), 'criterion' (set to None), 'epochs' (set to 5), 'scheduler' (set to None), 'device' (set to 'cuda'), and 'results' (set to False). The function is annotated with a docstring: '"""Trainer for PyTorch"""\n\n logger.info(f'Starting training...')\n training_start = time.time()\n\n train_losses, train_accs = [], []\n valid_losses, valid_accs = [], []\n\n model.to(device)'. The code uses standard Python syntax, including imports for 'logger' from 'logging', 'time' from 'time', and 'f' for formatted string literals.

Final data product: Modularized Neural Network classifiers





Thank You!

Questions?





Appendix





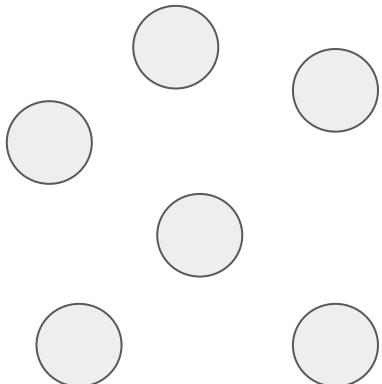
Appendix A: Edge Convolutions



Edge convolutions

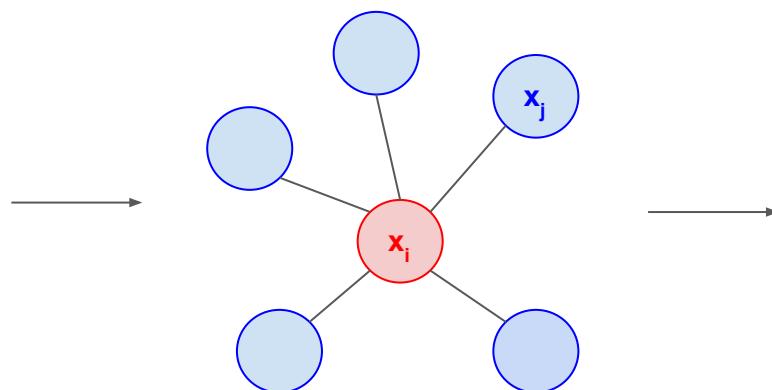
Step 1

Raw point cloud input



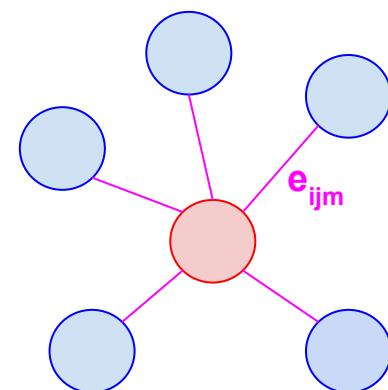
Step 2

Connect each node x_i to k nearest neighbors $x_j \dots x_k$ to create graph



Step 3

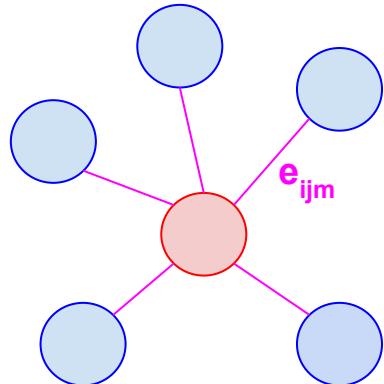
For each node and for each edge calculate edge features e_{ijm}



Edge convolutions continued...

Step 3

Calculate edge features for each edge e_{ijm}



Edge feature formula

- Weights θ and ϕ learned during training
- Encodes:
 - Global shape structure
 - Local information

$$e_{ijm} = \theta_m \cdot (x_j - x_i) + \phi_m \cdot x_i$$

Step 4

Apply activation and aggregation functions

Leaky ReLU
Activation
+
Max Pooling

Step 5 Output

Next EdgeConv layer
OR
Fully connected layer

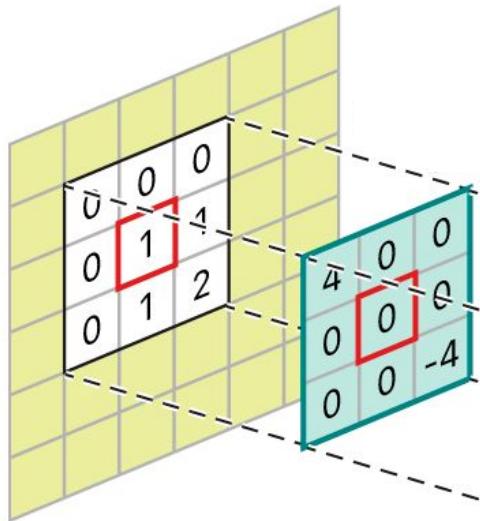


Appendix B: Image Convolutions vs. Edge Convolutions

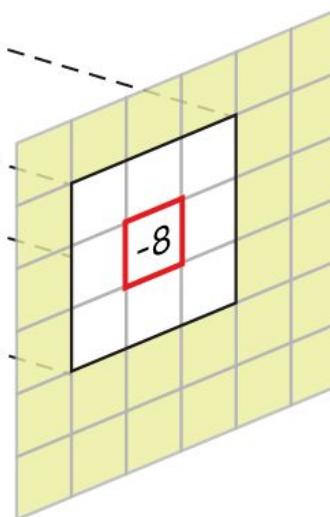


Image convolution: refresher

1) Input: image



2) Convolve filter (weights are learned)



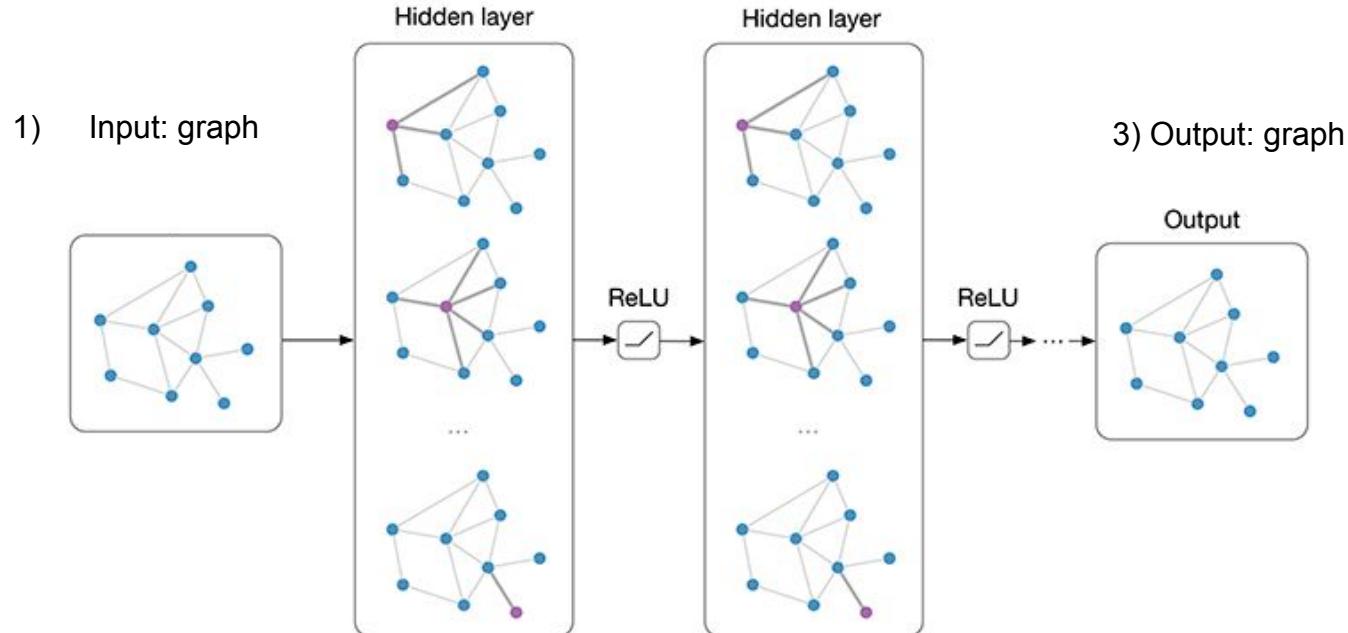
3) Output: New feature map

Then...

- Pooling
- Activation function

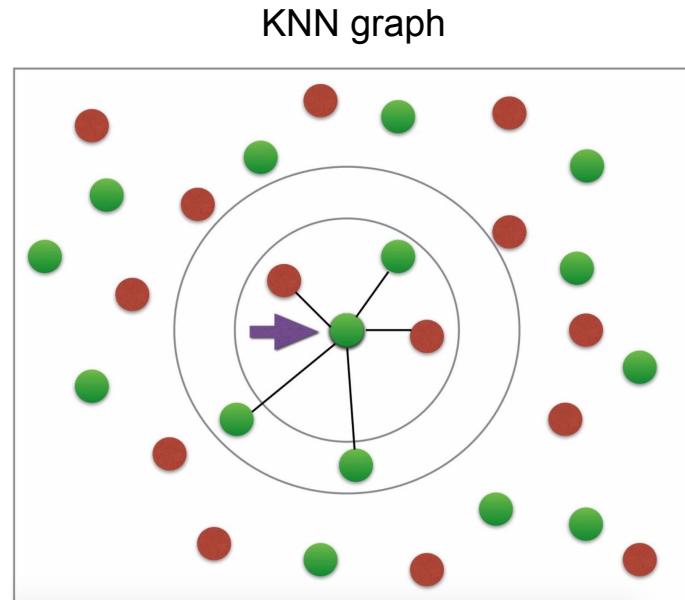
Traditional graph convolutions

2) Convolve filter and learn edge weights



Edge convolutions: improved graph convolutions

- Graph structure is not fixed!
- Graph is dynamically updated at each layer
- KNN algorithm is used
- Convolutions then performed on edges

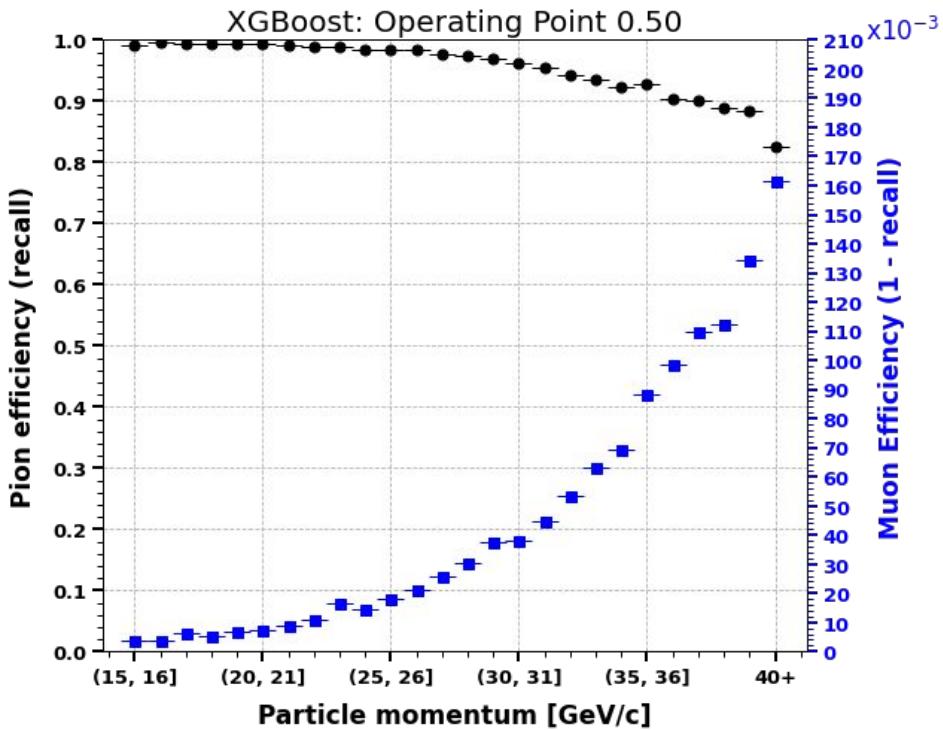
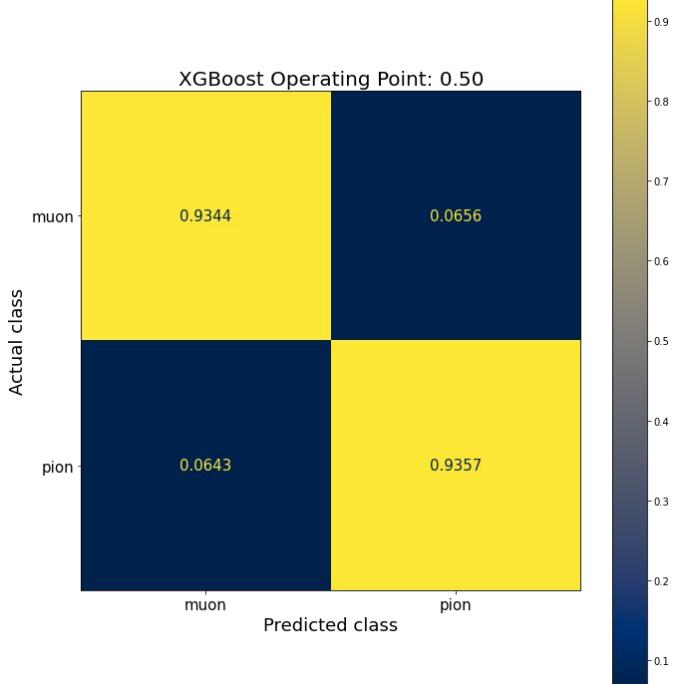




Appendix C: All Models with operating points 0.5

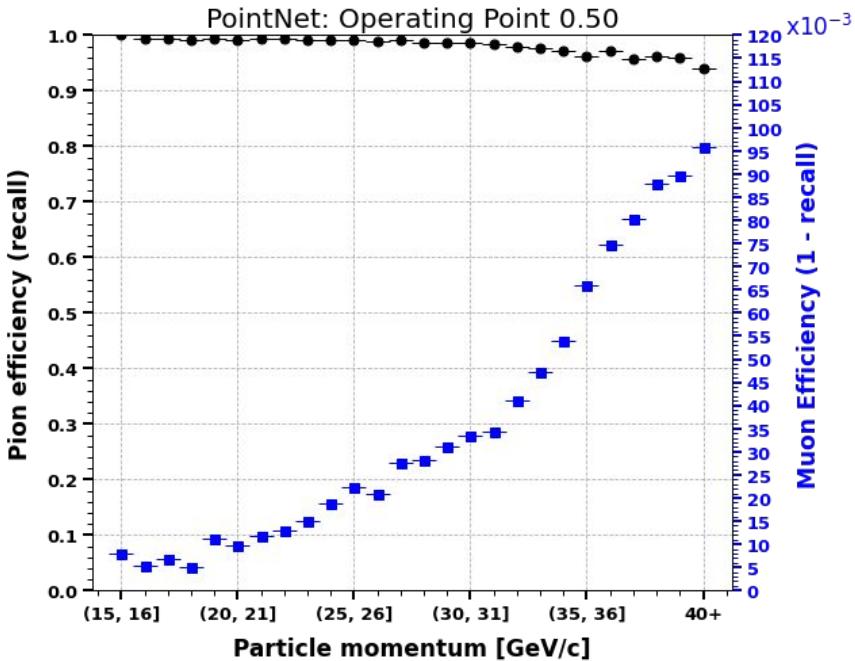
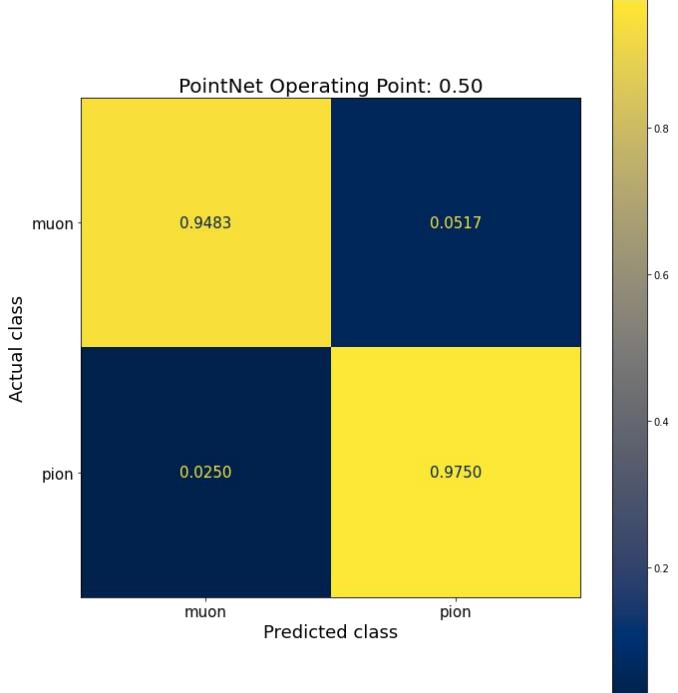


XGBoost with base 0.5 operating point



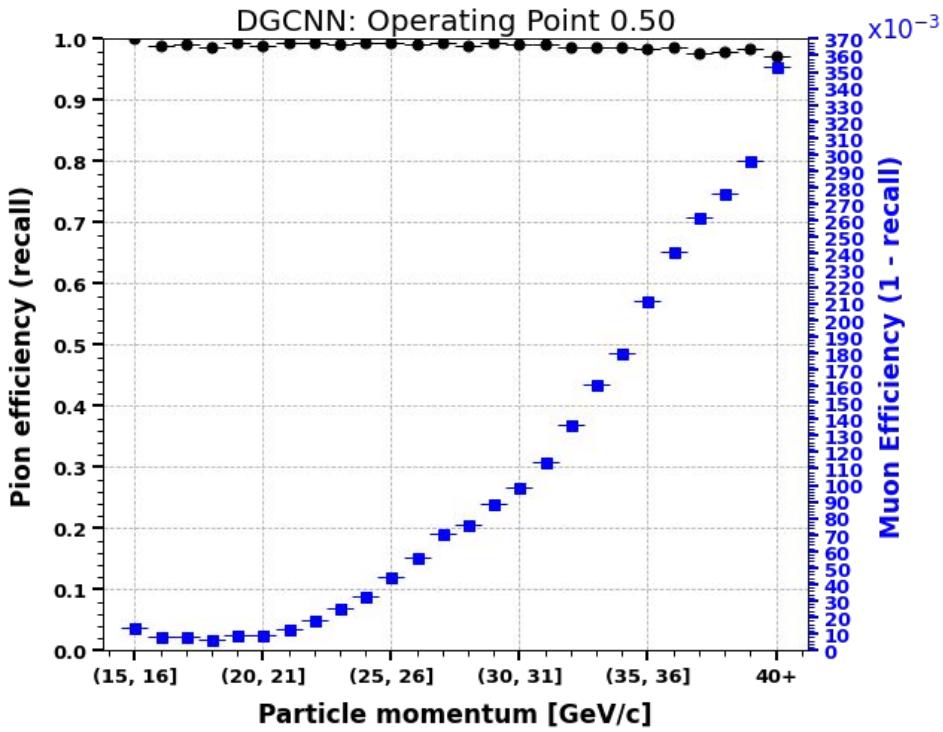
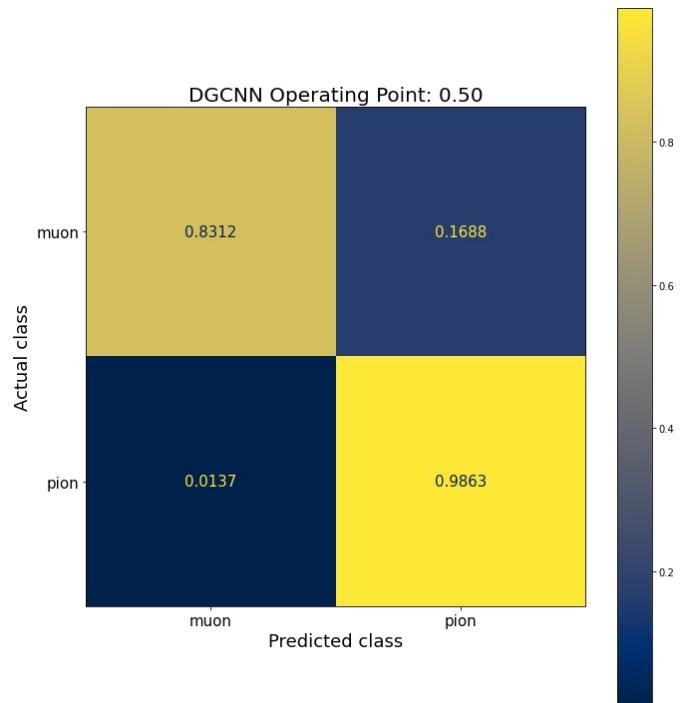
1. As a reminder, pion efficiency is recall for pion.

PointNet with base 0.5 operating point



1. As a reminder, pion efficiency is recall for pion.

Dynamic Graph CNN with base 0.5 operating point



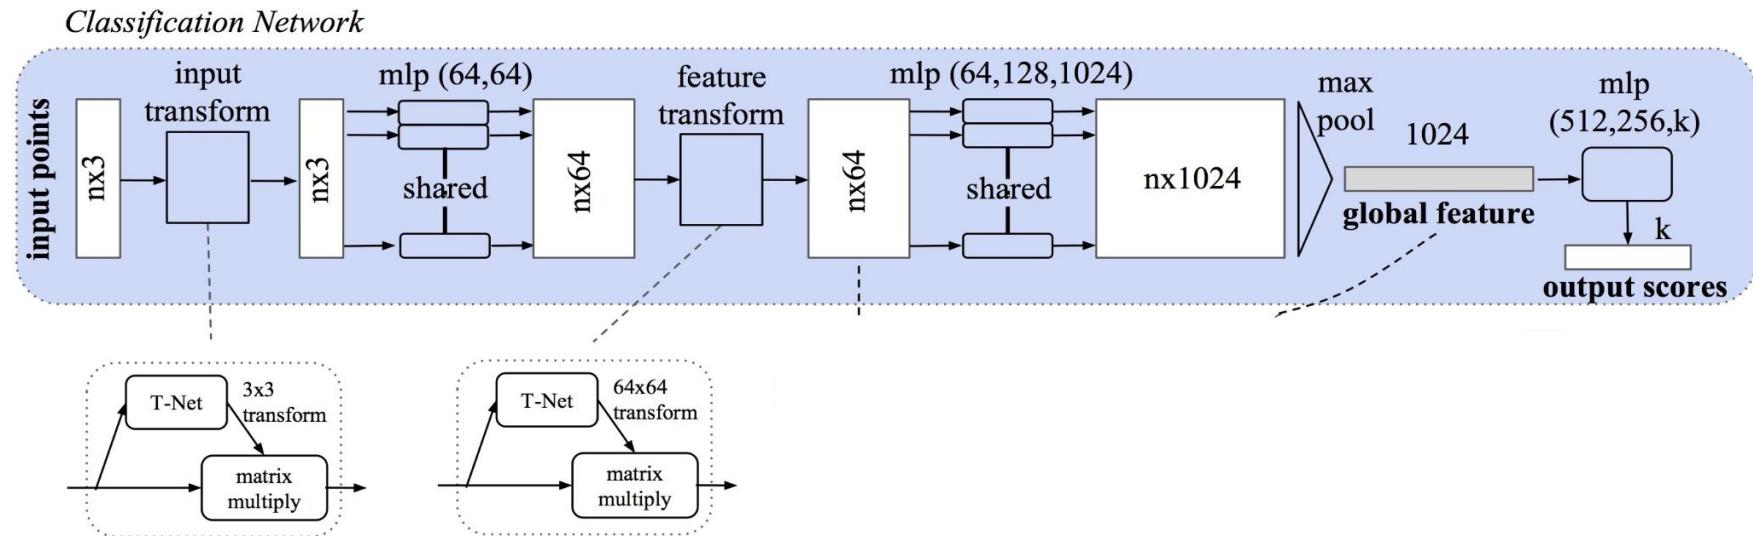
1. As a reminder, pion efficiency is recall for pion.



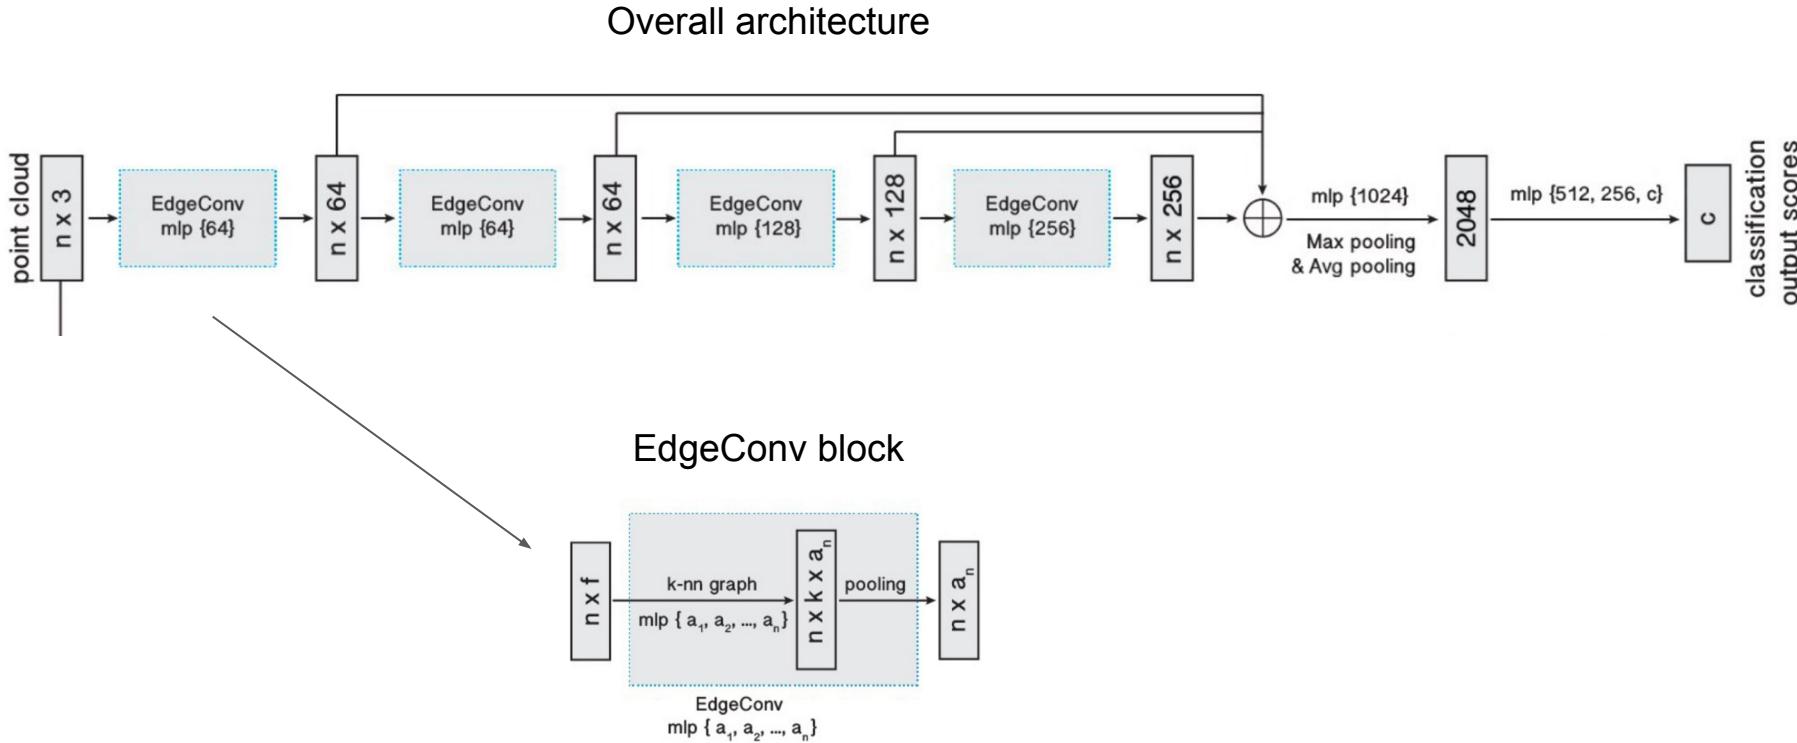
Appendix D: Deep learning architectures



PointNet Architecture



Dynamic Graph CNN Architecture

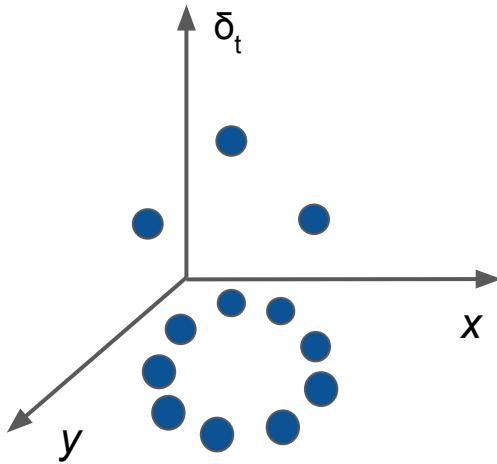




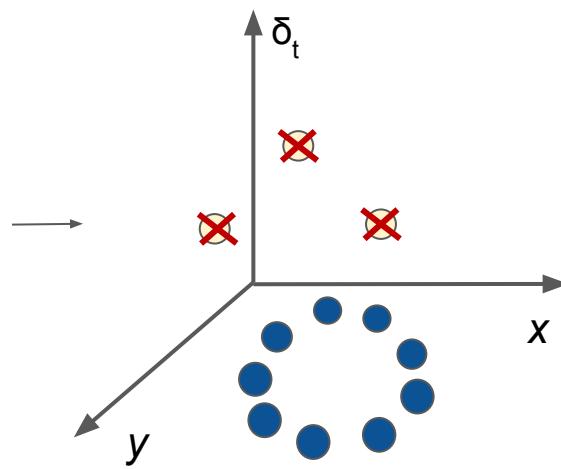
Appendix E: Physics Related



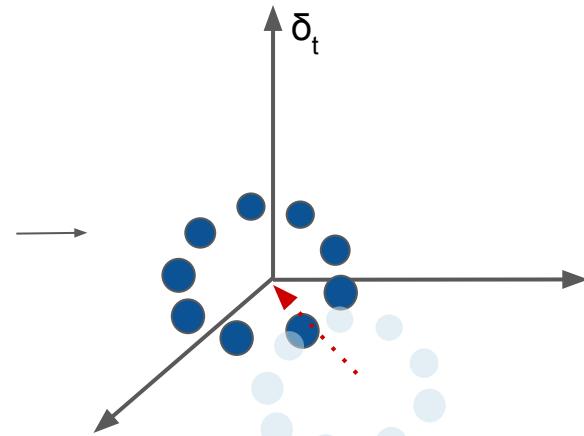
The feed into the deep learning models requires 2 filtering steps



Add time dimension to X,Y coordinates

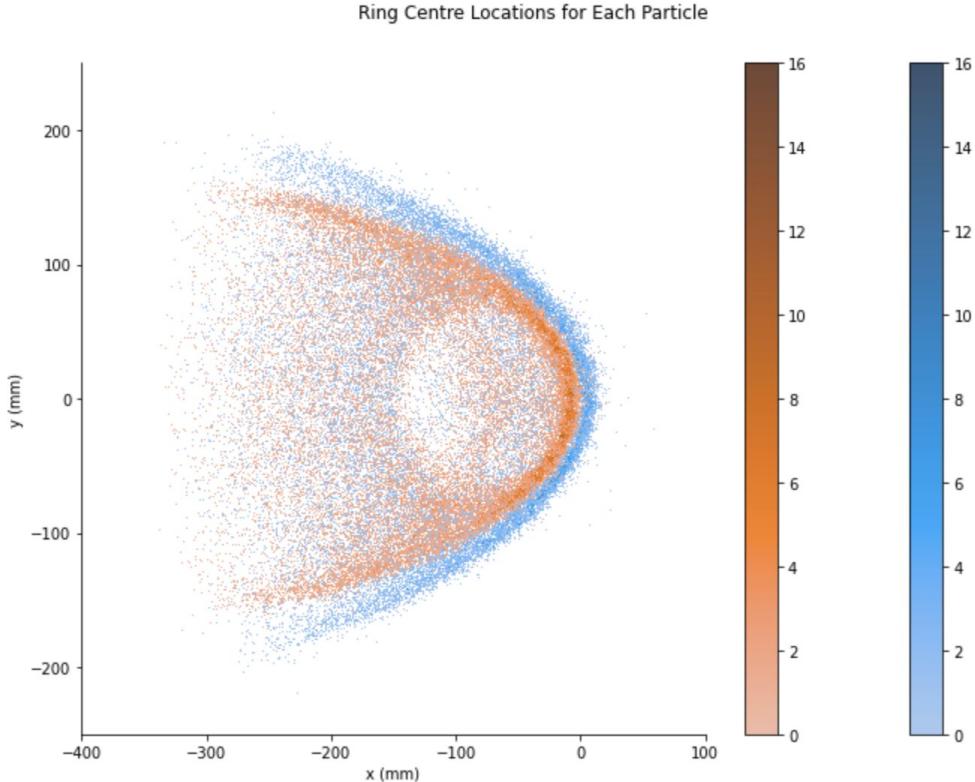


Filter for time



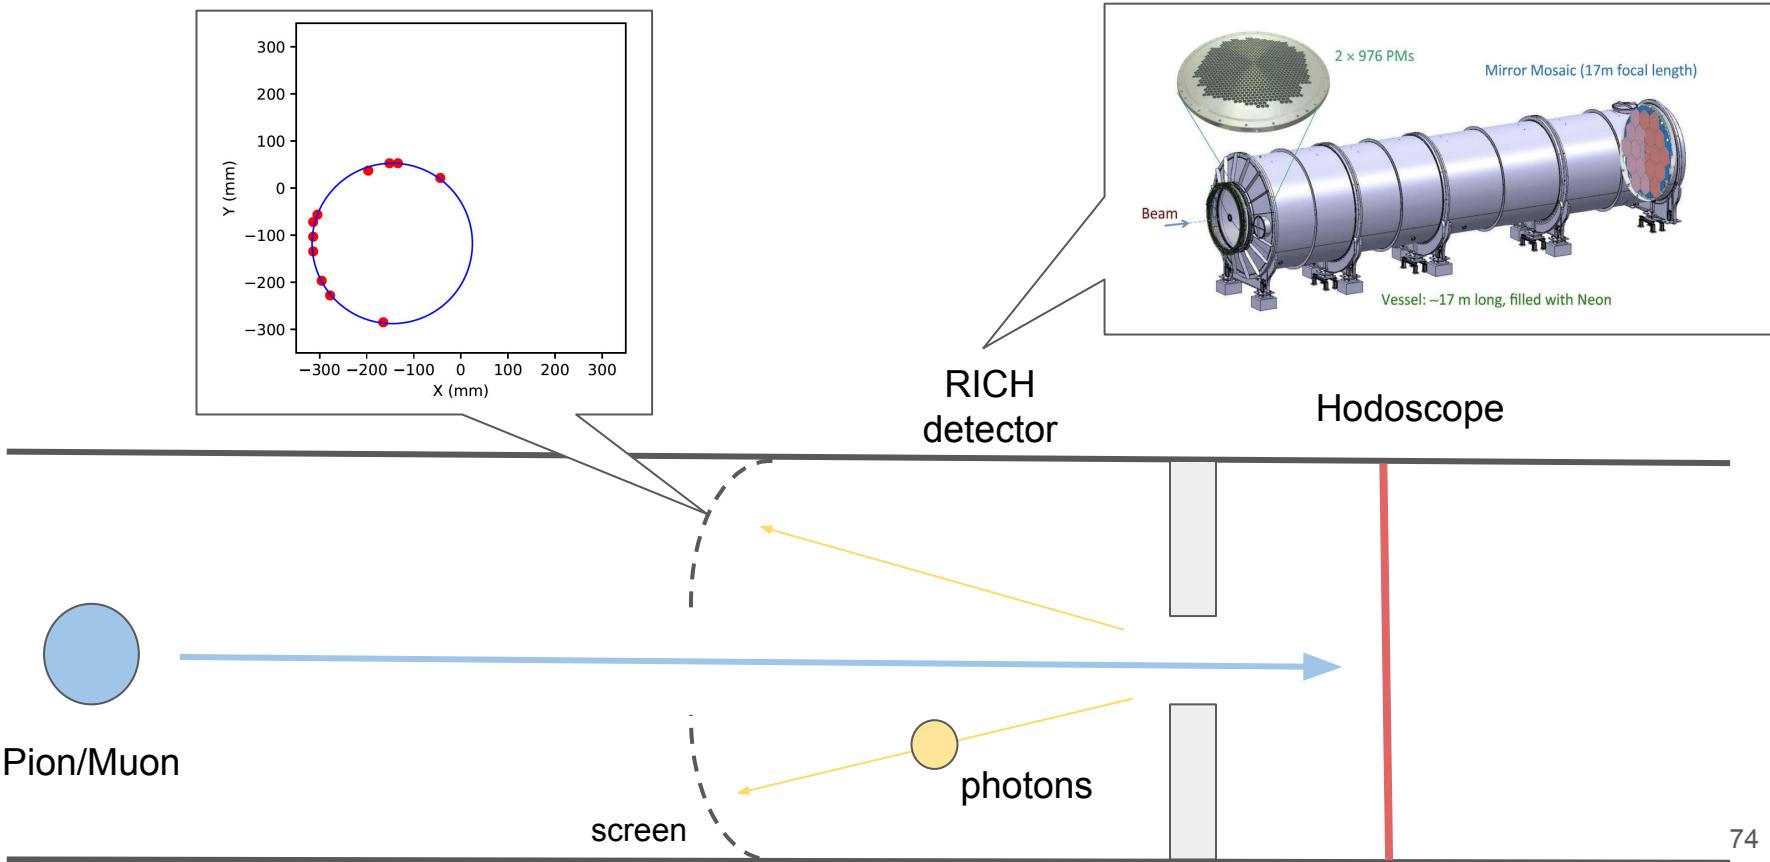
Recenter by
subtracting global
means for X and Y

There is a bias in the ring center

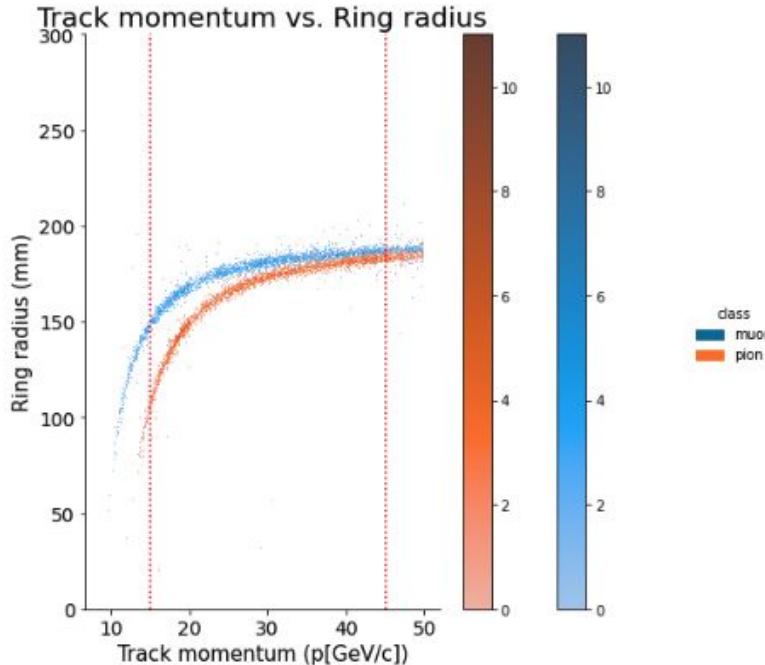


- Ring center distribution is **different** across the particles
- **Recentered** using the global mean X and Y positions

Each decay process produces light in the RICH detector



Undersampling is the answer but...



...need to account for the relationship between the momentum and ring radius

There is still a distinction in the ring radius between the particles

